

Implementation and Verification of a CAN-FD Bus Transceiver in VHDL

Mads Richardt (s224948)

February 28, 2026

Contents

Abstract	3
Abbreviations	3
1 Introduction	4
1.1 Motivation	4
1.2 Existing CAN Controller	4
1.2.1 Limitations	4
1.2.2 Decision to Redesign	5
1.3 Existing CAN FD IP Cores	5
1.3.1 Open-Source Implementations	5
1.3.2 Commercial Implementations	6
1.3.3 Rationale for In-House Development	6
1.4 Problem Statement	7
1.5 Objectives	7
2 Background	8
2.1 CAN as a Communication Bus	8
2.2 VHDL-2008 and OSVVM	9
3 Requirements	9
3.1 VHDL Code Standard and Design Constraints	9
3.1.1 Project-specific Infrastructure Requirements	9
3.1.2 File Structure and Naming	9
3.1.3 RTL Coding Style and Verification	10
3.2 From Specification to Structured Requirements	10
3.3 AI-Assisted Extraction: Utility and Limitations	12
4 CAN and CAN-FD Protocol Overview	12
4.1 Layered Reference Model	12
4.2 Frame Types and Formats	13
4.3 Bit Timing and Flexible Data Rate	14
4.4 Bit Stuffing	17
4.5 Cyclic Redundancy Check	17
4.6 Error Detection and Fault Confinement	17
5 Verification Plan	18
5.1 Layer	18
5.2 Side	19

5.3	Format Applicability	19
5.4	Observability	19
5.5	Priority	19
5.6	Requirement Distribution	20
5.7	Verification Plan Data Structure	20
5.7.1	Verification Method	21
5.7.2	Traceability: Label and File	21
5.7.3	Status	21
6	Design and Architecture	22
6.1	Ramifications of the Requirements Model on Initial Design Strategy	22
6.2	Architectural Design Decisions	23
6.2.1	Monolithic vs. Layered Architecture	23
6.2.2	Combined vs. Separated TX/RX FSMs	23
6.2.3	Per-Field vs. Per-Phase FSM Granularity	24
6.2.4	Internal LLC Frame Format	25
6.3	System Overview	25
6.4	LLC Frame Format	26
6.5	LLC Sub-layer	26
6.6	MAC Sub-layer	26
6.6.1	can_mac_fsm	26
6.6.2	can_mac_ser	28
6.6.3	can_mac_bs	28
6.6.4	can_mac_crc	28
6.7	FCE Sub-layer	28
6.8	PCS Sub-layer	29
7	Implementation	29
7.1	Interface Conventions	29
7.2	can_mac_fsm	30
7.2.1	FSM Structure and Mode Flag	30
7.2.2	TX Mode: Frame Transmission	32
7.2.3	RX Mode: Frame Reception	32
7.2.4	Error-Frame States	33
7.3	can_mac_ser	33
7.4	can_mac_bs	34
7.5	can_mac_crc	35
7.6	can_fce	36
7.7	can_pcs	37
7.7.1	Resynchronisation	37
7.7.2	Dual Bit Rate Switching	39
7.7.3	Transmitter Delay Compensation	39
8	Verification and Results	40
8.1	Testbench Results Summary	40
9	Discussion	41
9.1	Future Work	41
10	Conclusion	41
11	References	41
A	Verification Plan	41

A.1	Extracted Requirements	41
A.2	Verification Plan Fields	51

Abstract

*This document is a work-in-progress draft for the B.Eng thesis. Currently, the project is in the **Verification Plan** phase (Phase 2), with architectural prototyping for the transmitter sub-system completed.*

This thesis describes the design, implementation, and verification of a CAN (Controller Area Network) and CAN-FD (Flexible Data rate) transmitter sub-system. The implementation is targeting high-reliability engine controller applications and is compliant with the ISO 11898-1:2024 standard [1]. Key features include support for both Classic and FD frame formats, Transmitter Delay Compensation (TDC) for high-speed data phases, and a modular architecture separated into Link Layer Control (LLC), Media Access Control (MAC), and Physical Coding Sublayer (PCS).

Abbreviations

Table 1: Abbreviations used in this report.

Abbreviation	Meaning
ACK	Acknowledgement
AF	Acceptance Field
AUI	Attachment Unit Interface
CAN	Controller Area Network
CBFF	Classical Base Frame Format
CEFF	Classical Extended Frame Format
DF	Data Frame
DLL	Data Link Layer
EF	Error Frame
FBFF	FD Base Frame Format
FCE	Fault Confinement Entity
FEFF	FD Extended Frame Format
FTYP	Frame Type
LLC	Logical Link Control
LSDU	LLC Service Data Unit
MAC	Medium Access Control
OF	Overload Frame
OSI	Open Systems Interconnection
PCS	Physical Coding Sublayer
PDU	Protocol Data Unit
PMA	Physical Medium Attachment
RF	Remote Frame
SAP	Service Access Point
SDU	Service Data Unit
SOF	Start of Frame

Abbreviation	Meaning
TEC/REC	Transmit Error Counter / Receive Error Counter
VCID	Virtual CAN Channel Identifier

1 Introduction

TODO: Add a section on the IO extender board (The board is on the test wall, get the name from Alex)

1.1 Motivation

The Controller Area Network (CAN) has been the workhorse of automotive and industrial communication for decades. However, the increasing bandwidth requirements of modern systems led to the development of CAN-FD (Flexible Data rate), which allows for larger payloads and higher bit rates. This project aims to provide a robust, hardware-independent VHDL implementation of a CAN-FD transmitter.

1.2 Existing CAN Controller

The starting point for this project is an existing CAN Classic controller developed internally at the company. The controller is implemented in VHDL and has been integrated into a production IO-extender FPGA design. It supports CAN Classic frames with both 11-bit (base) and 29-bit (extended) identifiers at bit rates up to 500 kbit/s, and has been verified through hardware bring-up on physical CAN buses.

The controller follows a monolithic architecture where a single top-level wrapper (`can_bus_controller`) instantiates six sub-modules: a combined TX/RX frame FSM (`can_fsm`), a bit timing generator (`can_node_clock`), a dynamic bit stuffer (`can_stuff_bit_gen`), a CRC-15 engine (`gen_crc`), and two Avalon-ST converters for frame serialization and deserialization (`can_ast_to_serial`, `can_serial_to_ast`). The entire design totals roughly 2,000 lines of VHDL across seven source files.

The central component is `can_fsm`, an 810-line state machine with 18 states that handles both transmission and reception in a single process. The FSM manages frame arbitration, bit-level transmission and reception, stuff bit error checking, CRC validation, ACK handling, and error flag generation. Error counting (TEC/REC) and node state transitions (Error Active, Error Passive, Bus Off) are handled in separate processes within the same file but are tightly coupled to the main FSM through shared signals.

1.2.1 Limitations

While the existing controller is functional for CAN Classic, several architectural limitations prevent it from being extended to support CAN FD:

Single bit rate domain. The `can_node_clock` module generates a single pair of timing strobes - one sample pulse and one transmit pulse - derived from a fixed set of bit timing parameters. CAN FD requires switching between a nominal bit rate (used during arbitration) and a faster data bit rate (used during the data phase), with Transmitter Delay Compensation (TDC) to account for the transceiver round-trip delay at the higher rate. Adding dual bit rate support and TDC to the existing clock module would require a fundamental redesign of the timing architecture.

Dynamic bit stuffing only. The `can_stuff_bit_gen` module implements the CAN Classic rule of inserting an inverse bit after five consecutive identical bits. CAN FD introduces a second stuffing mode - fixed bit stuffing - where a stuff bit is inserted at fixed intervals during the CRC field, and a Stuff Bit Count (SBC) field with Gray-coded parity is appended. The existing stuffer has no mechanism for mode switching or SBC generation.

Single CRC polynomial. The controller uses a single CRC-15 instance. CAN FD requires three CRC polynomials: CRC-15 for Classic frames, CRC-17 for FD frames with payloads up to 16 bytes, and CRC-21 for larger FD payloads. Furthermore, the CRC data feed differs between Classic and FD - in FD frames, dynamic stuff bits in the arbitration region are included in the CRC computation, requiring a dual data feed to the CRC engine.

Combined TX/RX FSM. The monolithic FSM interleaves transmission and reception logic in every state, with the `is_transmitter` flag selecting the active code path. This coupling makes it difficult to add FD-specific states (such as BRS, ESI, and the SBC field) without increasing the already high cyclomatic complexity. A CAN FD frame has more control fields than a Classic frame, and handling both TX and RX paths for all six frame variants (CB, CE, FB, FE data frames, plus remote frames for CB and CE) in a single process would result in an unwieldy state machine.

Embedded error handling. Error detection, error flag transmission, and error counter management are distributed across the main FSM process and its auxiliary processes. The ISO 11898-1 standard defines the Fault Confinement Entity (FCE) as a logically separate component with well-defined interfaces to the MAC and PCS sub-layers. Extracting the error handling into a reusable, independently testable FCE module - as required by the standard's layered architecture - would require significant refactoring of the existing FSM.

1.2.2 Decision to Redesign

Given these limitations, extending the existing controller to CAN FD would require modifying nearly every sub-module and fundamentally restructuring the FSM. The resulting design would carry the constraints of the original monolithic architecture while trying to support a significantly more complex protocol. Instead, a clean-sheet redesign was chosen, structured around the ISO 11898-1 layered architecture (LLC, MAC, PCS, FCE) with separated TX and RX paths, reusable sub-components (bit stuffer, CRC engine), and typed record interfaces. This approach allows each sub-layer to be implemented and verified independently, supports both Classic and FD frame formats from the outset, and produces a modular design that can be maintained and extended without cross-cutting changes.

1.3 Existing CAN FD IP Cores

Before committing to an in-house redesign, the available CAN FD controller IP cores were evaluated. Table 2 summarizes the candidates, spanning both open-source and commercial offerings.

1.3.1 Open-Source Implementations

CTU CAN FD [2], [3] is the only mature open-source CAN FD controller available as synthesizable HDL. Developed at the Czech Technical University in Prague, it is written in VHDL, licensed under MIT, and has been conformance-tested against ISO 16845-1 [4]. The controller includes a full TX and RX pipeline with up to four TX buffers, acceptance filtering, timestamping, and a register interface with DMA support. A mainline Linux kernel driver has been available since kernel version 5.12. CTU CAN FD represents a complete, production-oriented CAN node - a significantly broader scope than what is needed in this project.

OpenCores CAN [5] is a Verilog controller modeled after the Philips SJA1000 register interface. It is one of the earliest open-source CAN cores and is widely cited in academic work. However, it supports only CAN 2.0B (Classic CAN) and has seen no active development since approximately 2010. It cannot serve as a starting point for CAN FD.

Canola [6] is a VHDL CAN 2.0B controller with a clean VHDL-2008 codebase and a cocotb-based test-bench. It includes a Triple Modular Redundancy (TMR) wrapper for radiation-tolerant applications. Like the OpenCores core, it does not support CAN FD.

1.3.2 Commercial Implementations

Bosch M_CAN [7], [8] is the reference CAN FD controller, developed by the inventor of both CAN and CAN FD. M_CAN is the IP core embedded in virtually every automotive microcontroller (NXP S32, Infineon AURIX, STM32, TI Jacinto, Renesas RH850). It is licensed under NDA with per-design royalty fees.

AMD/Xilinx CAN FD [9] is a soft IP core included in the Vivado Design Suite. It provides an AXI4-Lite register interface with up to 32 acceptance filters, TX mailboxes, and RX FIFOs. It is device-locked to AMD/Xilinx FPGAs and cannot be ported to other targets.

CAST CAN FD [10] is a technology-independent RTL core with AMBA APB/AHB bus interface options. It is licensed per-design with an upfront fee. Synopsys (DesignWare) and Cadence offer similar ASIC-targeted CAN FD cores under their respective IP licensing programs.

Table 2: Survey of available CAN FD controller IP cores.

Implementation	Language	CAN FD	License	Scope	Conformance Tested
CTU CAN FD [2]	VHDL	Yes	MIT	Full node (TX+RX, buffers, DMA)	ISO 16845-1
OpenCores CAN [5]	Verilog	No	LGPL	Full node (CAN 2.0B only)	No
Canola [6]	VHDL	No	MIT	Full node (CAN 2.0B, TMR)	No
Bosch M_CAN [7]	HDL (NDA)	Yes	Per-design royalty	Full node	Yes (reference)
AMD/Xilinx CAN FD [9]	HDL	Yes	Vivado-included	Full node	Yes
CAST CAN FD [10]	RTL	Yes	Per-design fee	Full node	Yes

1.3.3 Rationale for In-House Development

Despite the availability of these solutions, none satisfies the combined requirements of a large engine design company developing safety-critical control systems. The decision to develop the CAN FD transceiver in-house was driven by five considerations.

TODO: Add something about the 30 year maintenance requirements that the company is responsible for (hard to rely on 3rd party products for this long). **IP ownership and supply chain independence.** Commercial IP cores introduce a licensing dependency on an external vendor. For a product with a multi-decade

lifecycle - typical in heavy-duty engine applications - this dependency carries risk: vendors may discontinue support, change licensing terms, or be acquired. Owning the RTL outright eliminates this exposure and ensures that the design can be maintained, ported, and modified without third-party approval for the full product lifetime. The open-source CTU CAN FD avoids the licensing risk, but using it still means adopting a codebase whose architecture, naming conventions, and design decisions were made for a different context.

Verification authority. In safety-critical domains, the verification evidence must be traceable from standard requirements to RTL assertions and testbench results. Adopting a third-party core - even one conformance-tested against ISO 16845-1 - means inheriting its verification artifacts rather than producing them. The company's verification methodology requires full control over the verification plan, the testbench architecture, and the assertion coverage. Building the RTL in-house allows the verification plan (described in Section ??) to drive the implementation, ensuring that every module is verified against the specific requirements extracted from the standard, using the company's own toolchain and conventions.

Architectural scope. All available IP cores implement a complete CAN node: TX and RX pipelines, message RAM, acceptance filtering, buffer management, register interfaces, and in some cases DMA controllers. The company's application requires only the data link layer (LLC, MAC, FCE) and physical coding sublayer (PCS) - the protocol engine that converts between byte-level frame data and the serial bus. The higher-level buffering and filtering logic already exists in the company's FPGA infrastructure. Adopting a full-node IP core would introduce unnecessary complexity and area overhead, and the integration effort to bypass or disable the unused subsystems may approach the effort of a targeted implementation.

Integration with existing infrastructure. The company's FPGA designs use a specific Avalon-ST streaming interface for inter-module communication, a particular clock and reset architecture, and established conventions for signal naming and module boundaries. A third-party core would require an adaptation layer to bridge its native interface (AXI, APB, or custom register map) to the existing infrastructure. The in-house design uses the company's interface conventions natively, eliminating this integration overhead.

Platform independence. The AMD/Xilinx CAN FD core is locked to Xilinx devices. The Bosch M_CAN and other commercial cores are delivered as technology-specific netlists or encrypted RTL for a particular target. The in-house design is written in portable VHDL-2008, synthesizable on any FPGA platform or ASIC flow, ensuring that the IP remains usable if the company changes FPGA vendors.

1.4 Problem Statement

The need for CAN FD support in the company's engine controller platform, combined with the architectural limitations of the existing CAN Classic controller (Section 1.2.1) and the unsuitability of available third-party IP cores for the company's specific requirements (Section 1.3.3), motivates the development of a new CAN FD transceiver from the ground up. The challenge lies in creating a modular, independently verifiable implementation that supports both Classic and FD frame formats, handles dual bit rate switching with Transmitter Delay Compensation (TDC), and maintains strict compliance with ISO 11898-1 [1] - all while fitting into the company's existing FPGA infrastructure and verification methodology.

1.5 Objectives

- Implement a VHDL-2008 compliant CAN/CAN-FD transmitter.
- Support both Base (11-bit) and Extended (29-bit) identifiers.
- Implement TDC measurement and compensation logic.

- Ensure high verification coverage using OSVVM and GHDL.

2 Background

2.1 CAN as a Communication Bus

The Controller Area Network (CAN) is a serial communication bus developed by Bosch in 1986 ? to connect electronic control units (ECUs) in automotive environments without a central host computer. Where point-to-point wiring and star-switched architectures require a dedicated conductor between every communicating pair, CAN uses a shared two-wire differential bus on which all nodes broadcast simultaneously and arbitrate access without any designated bus master. Any node may initiate a transmission at any time; contention is resolved by a non-destructive bitwise arbitration in which the transmitter with the lower-priority identifier detects the collision and silently withdraws, leaving the winner's frame intact. Differential signaling on a twisted pair (ISO 11898-2 physical layer) provides strong common-mode noise rejection - a practical necessity in the electrically harsh environment of an engine bay or industrial cabinet.




Figure 1: CAN bus consisting of three CAN nodes connected via a shared differential two-wire bus. Each node contains a CAN controller and transceiver; termination resistors at each end of the bus prevent signal reflections.

CAN's error-handling architecture is a distinguishing feature relative to simpler serial protocols. Five complementary error detection mechanisms operate concurrently on every transmitted frame: bit monitoring, frame check, cyclic redundancy checking, acknowledgement checking, and bit stuffing violation detection. A fault confinement mechanism tracks each node's error history and automatically escalates from error-active operation through error-passive to a bus-off state in which a persistently faulty node is electrically isolated from the network, preventing it from corrupting communication between healthy nodes. Together these properties made CAN the protocol of choice for safety-relevant in-vehicle networks; adoption subsequently spread to industrial automation, medical devices, and aerospace ground support equipment.

CAN's original data payload was capped at eight bytes per frame, limiting raw throughput to around 1 Mbit/s. As embedded control applications became more data-intensive, this ceiling became a practical

constraint. CAN FD (Flexible Data-Rate), finalized by Bosch in 2012 and incorporated into ISO 11898-1 in 2015, extends the maximum payload to 64 bytes and introduces a separate higher-speed data phase with bit rates up to 8 Mbit/s or beyond, while preserving the Classic CAN arbitration phase and the fault-confinement architecture unchanged. The governing standard for this project is ISO 11898-1:2015, which specifies both Classic CAN and CAN FD data link layer and physical signaling requirements.

2.2 VHDL-2008 and OSVVM

The implementation language for this project is VHDL-2008, the current revision of the IEEE VHDL standard. VHDL-2008 adds several features relevant to parameterized hardware design and verification: unconstrained record elements, enhanced generic lists, and improved support for the IEEE numeric packages. The GHDL open-source simulator [11] supports VHDL-2008 natively and is used for all simulation in this project.

The verification framework is OSVVM (Open Source VHDL Verification Methodology) [12], a VHDL-native library providing test infrastructure including clock and reset generation, constrained-random stimulus, functional coverage, and a uniform pass/fail reporting framework. OSVVM procedures replace ad-hoc signal manipulation in testbenches, ensuring that timing relationships are expressed in terms of clock cycles rather than time literals and that pass/fail decisions are logged uniformly across all test scenarios.

3 Requirements

Bridging the gap between a normative specification document and a structured, verifiable requirements set is a pivotal task in any protocol implementation project. The present section describes the process of distilling a coherent set of verifiable requirements from the ISO 11898-1 standard. In addition, the constraints imposed by general company practices and standards are described along with project-specific requirements related to the larger system in which the module is intended to integrate.

3.1 VHDL Code Standard and Design Constraints

The constraints on this project come from two distinct sources. Two requirements are specific to this project's place within the company's existing CAN infrastructure while the remaining constraints come from the company's VHDL Code Standard and apply uniformly to all FPGA IP modules developed in-house.

3.1.1 Project-specific Infrastructure Requirements

1. **Avalon-ST user interface:** The CAN controller's external interface to the host system must use the Avalon-ST streaming protocol [13] (data, valid, ready, sop, eop). The requirement applies specifically to the boundary between the CAN controller and its user.
2. **IP library CRC block:** The company maintains a reusable, parameterised CRC generator (`gen_crc`) in its IP library. This module must be used for all CRC computations.

3.1.2 File Structure and Naming

1. **Per-module file structure:** Each module must be organized into a fixed directory layout: `hdl_src/` for RTL source files, `hdl_tb/` for test bench files, and `test_case/` for waveform configuration. One entity per VHDL file, with the filename matching the entity name. Every entity requires a dedicated test bench, unless it is instantiated exclusively as a sub-module of a fully tested parent.

2. **Entity port types:** Entity ports are restricted to `std_logic`, `std_logic_vector`, and records or arrays of these types. Modes are restricted to `in` and `out`.
3. **Naming conventions:** A mandatory prefix/suffix scheme applies to all VHDL identifiers: types (`t_`), constants (`c_`), generics (`gc_`), processes (`p_`), functions (`f_`), packages (`pk_`), state variables (`s_`), entity inputs (`_i`), and entity outputs (`_o`).

3.1.3 RTL Coding Style and Verification

1. **RTL design rules:** Synchronous processes must be sensitive to the clock only. Reset must be synchronous and initialize all control registers. FSMs are preferably implemented as single-process designs where all signal assignments are derived from the current state, with an explicit other-state that returns to a known safe state.
2. **Test bench requirements:** Test benches must follow a black-box testing model and test cases must be ordered as: reset tests first, then a normal-usage test, then all remaining tests.

3.2 From Specification to Structured Requirements

The requirements engineering process was aimed at tackling two key objectives:

1. Extracting a clear and actionable set of requirements that could serve as a starting point for the design phase.
2. Establishing a clear, traceable link between the ISO 11898-1 specification and the verification environment.

Both objectives are complicated by the nature of the source material. Normative requirements are distributed across subsections, often restated from different perspectives, and interspersed with explanatory and descriptive text. Accordingly, extracting a precise unambiguous set of requirements from such prose is non-trivial and inherently prone to oversights and misinterpretations. Nonetheless, a structured and precise requirements set is absolutely essential. It serves as both the starting point for the system design phase and as the target of verification effort. [14]

The requirements set was constructed using the AI-assisted pipeline shown in Figure 2. The first step was converting the ISO 11898-1 pdf to Markdown - a format which can be efficiently searched and ingested by LLM models. The resulting Markdown file was then fed to a Claude Sonnet 4.6 LLM agent, which was prompted to extract all normative statements - sentences containing words like “shall”, “should”, “must”, and their corresponding negations.

This process yielded a raw set of 168 normative statements linked to the ISO standard sections from which they were extracted. The normative set was then manually reviewed, consolidated, and distilled into a final set of 38 requirements (reproduced in Section A).

The requirements set is stored as a TOML file (`verification_plan.toml`), one `[[requirement]]` block per entry. To support ongoing AI-assisted refinement without the risk of silent data corruption, a custom Model Context Protocol (MCP) server (`verification_plan_manager.py`) was developed alongside the plan. The server exposes query, update, insert, and statistics operations as tool calls that the AI coding agent can invoke directly within its development environment. Each write operation targets a single requirement entry and validates field values against the schema before committing. This makes it structurally impossible for the agent to silently drop entries, fabricate field values, or corrupt TOML syntax - failure modes that arise inevitably when an LLM is asked to rewrite a large structured file in one operation. During the requirements phase, the agent worked exclusively with the five fields relevant at this stage:




Figure 2: Pipeline generating the `verification_plan.toml` artifact. 1) LLM agent extraction of normative statements from the ISO 11898-1 standard. 2) Manual grouping of related normative statements and requirement distillation. 3) Augmentation with additional requirement labels and verification specific fields.

- **source_clause:** The ISO 11898-1 section reference (e.g. §6.6.13.1). This is the traceability anchor - every requirement links back to the clause from which it was distilled, making it possible to verify the requirements set against the standard during review.
- **original_wording:** Verbatim ISO text for the relevant clauses. Preserving the source wording prevents paraphrase drift and provides a fallback for resolving ambiguity during implementation.
- **paraphrase:** A concise, implementer-facing restatement of the requirement. Where the ISO prose bundles multiple obligations into a single clause, the paraphrase enumerates them as numbered sub-claims, each independently verifiable.
- **priority:** A three-level rating - P1 (need-to-have, derived from “shall” obligations and core correctness), P2 (verified in the normal cycle), or P3 (optional, derived from “should” clauses or implementation-dependent features). Priority drives implementation sequencing and scope decisions when schedule is constrained. Of the 38 requirements, 31 are rated P1, 5 are P2, and 2 are P3. Every requirement not rated P1 was explicitly demoted; the rationale for each demotion is given in Table 3.
- **notes:** Residual clarifications not resolved by the paraphrase - implementation constraints, out-of-scope markers, or known ambiguities flagged for design review.

Table 3: Priority demotion rationale for all requirements not rated P1.

ID	Topic	Priority	Demotion rationale
REQ-002	LLC TX request and abort timing	P2	The 2-SOF processing window is a responsiveness guarantee, not a correctness constraint. A node that transmits eventually but outside this window sends valid frames.
REQ-011	Remote frame	P2	A data-only node is a valid CAN implementation. Remote frame support is a distinct feature subset not required for basic interoperability.
REQ-016	ESI bit transmission	P2	ESI communicates the node’s error state as an informational signal. Incorrect ESI does not abort a frame or trigger a protocol error at any receiver.

ID	Topic	Priority	Demotion rationale
REQ-036	MAC data consistency	P2	A frame corrupted mid-transmission would be error-flagged by other nodes on the bus and retransmitted. The requirement reduces bus pollution from invalid frames but the node remains functional without it.
REQ-037	Error signalling enable	P2	Error signalling itself is covered by P1 requirements. This requirement concerns only the existence of a configurable disable mode, which is an optional operational feature.
REQ-004	Frame acceptance filtering	P3	ISO uses advisory “should” language. Acceptance filtering is also an application-layer concern above the LLC boundary verified here.
REQ-038	DLC padding	P3	Padding with 0xCC applies only when the implementation exposes a configurable maximum-data-byte restriction. The feature may be waived entirely if that restriction is not implemented.

3.3 AI-Assisted Extraction: Utility and Limitations

The LLM agent earned its keep by bootstrapping and linking the initial normative statement set. Having a fully populated and linked starting point - even one requiring substantial revision - gave the manual review process a concrete artifact to work from. That said, the overall time saving was likely marginal. The agent’s output had to be reviewed statement by statement, which is functionally similar to extracting requirements manually in the first place. The primary benefit of the AI-assisted approach was therefore not efficiency, but rather the increased consistency associated with an automated extraction process.

The 38 requirements, each linked to its ISO source clause and assigned a priority, define the scope of what must be implemented and verified. They also function as a structured map to the protocol: every requirement points to a mechanism that must be understood before implementation can begin. The following section provides that understanding - covering the sub-layer model, frame formats, bit timing, bit stuffing, CRC, and error handling - so that the verification plan that follows can be grounded in full protocol context.

4 CAN and CAN-FD Protocol Overview

This section covers the ISO 11898-1 layered reference model, frame types and fields, bit timing and the dual-rate mechanism that distinguishes CAN FD from Classic CAN, bit stuffing, CRC, and error handling. Each mechanism is cross-referenced to the corresponding verification plan entries (REQ-NNN) established in Section 3. A reader already familiar with ISO 11898-1 may skip to Section 5.

4.1 Layered Reference Model

ISO 11898-1 structures the CAN data link layer into three functional sub-layers and a cross-cutting Fault Confinement Entity (FCE):

- **LLC (Logical Link Control):** accepts frame requests from the host application, applies retransmission policy on error or lost arbitration, and supplies frames to the MAC in serialized form.
- **MAC (Medium Access Control):** encodes and decodes the frame bit-by-bit - performing bit stuffing and destuffing, CRC generation and checking, and acknowledgement handling - and governs bus access arbitration.




Figure 3: ISO 11898-1 CAN node reference model. The data link layer comprises three functional sub-layers - LLC, MAC, and PCS - and a cross-cutting Fault Confinement Entity (FCE). The LLC accepts service requests from the host application; the MAC encodes and decodes frames at the bit level, performing bit stuffing, CRC, and acknowledgement handling; the PCS generates sample-point timing and interfaces to the Physical Medium Attachment (PMA); the FCE maintains error counters and governs node-state transitions. Each sub-layer maps to a dedicated VHDL module in this implementation.

- **PCS (Physical Coding Sub-layer):** manages bit timing, clock synchronization (including Transmitter Delay Compensation for FD data phase), and the sample/drive interface to the physical transceiver.
- **FCE (Fault Confinement Entity):** maintains Transmit Error Counter (TEC) and Receive Error Counter (REC), escalating the node's error state from Error Active through Error Passive to Bus Off as error counts accumulate.

In the implementation described in this report, each sub-layer maps to a dedicated VHDL module, and the sub-layer interfaces become the port records connecting those modules (Section 6). LLC service obligations are captured in [REQ-001](#) through [REQ-005](#) and [REQ-033](#); MAC frame-encoding rules in [REQ-006](#) through [REQ-024](#) and [REQ-032](#) through [REQ-038](#); PCS timing constraints in [REQ-025](#) through [REQ-028](#); and FCE counter and state-transition rules in [REQ-029](#) through [REQ-031](#) and [REQ-037](#).

4.2 Frame Types and Formats

CAN defines two classes of frames: Classic CAN (CC) and CAN FD (FD). Within each class, frames may carry either an 11-bit base identifier or a 29-bit extended identifier, giving four frame formats: CB (Classic Base), CE (Classic Extended), FB (FD Base), and FE (FD Extended), as shown in Figure 4. Classic frames (CB and CE) additionally support remote frame variants (RTR=1, no data field), giving six bus frame types in total. CAN XL frames are out of scope for this project.

A Classic CAN frame consists of Start of Frame (SOF), Arbitration field (identifier, RTR, IDE), Control field (DLC), Data field, CRC field, ACK slot and delimiter, End of Frame (EOF), and Intermission. SOF is a single dominant bit that marks the beginning of a frame and triggers hard synchronization in all receiving nodes ([REQ-010](#)). Within the arbitration field, bits are transmitted MSB first ([REQ-034](#)); the RTR bit distinguishes data frames from remote frames, being dominant for data frames and recessive for remote frames ([REQ-011](#)); and in extended frames an SRR placeholder bit transmitted recessive precedes the IDE bit ([REQ-012](#)). The DLC encodes the number of data bytes using the four-bit mapping defined in ISO 11898-1 ([REQ-033](#), [REQ-038](#)). After the data and CRC fields, the ACK slot carries a dominant bit driven by every receiver that has successfully validated the frame CRC - the transmitter monitors this slot and reports an acknowledgement error if no dominant bit is received ([REQ-014](#)). The frame is delimited by seven recessive EOF bits followed by three recessive intermission bits ([REQ-020](#), [REQ-008](#)).

A CAN FD frame shares the same structure through the arbitration phase and then introduces FD-specific control fields. The FDF bit distinguishes an FD frame from a Classic frame; a recessive FDF triggers the FD control field sequence including reserved bits, BRS, and ESI ([REQ-015](#)). The BRS (Bit Rate Switch) bit

controls the transition to the data-phase bit rate: when BRS is recessive the bus switches to the faster data rate immediately after the BRS sample point and returns to the nominal rate at the CRC delimiter (REQ-032). The ESI (Error State Indicator) bit reflects the transmitting node's fault-confinement state: a node in Error Passive state shall transmit ESI recessive (REQ-016).

4.3 Bit Timing and Flexible Data Rate

Every CAN bit period is divided into four non-overlapping time segments measured in Time Quanta (TQ), where one TQ equals the period of the prescaled system clock. CAN FD extends Classic CAN with an independently configured data-phase bit rate; a CAN FD node therefore maintains two independent sets of segment parameters, one for the nominal rate and one for the data rate (REQ-025):

- **Sync Segment (SYNC_SEG)**: one TQ; the point at which the bus is expected to produce a recessive-to-dominant edge after synchronization.
- **Propagation Segment (PROP_SEG)**: compensates for round-trip signal propagation delay on the bus and in the transceiver. It shall be programmed to be at least as long as twice the maximum bus propagation delay.
- **Phase Segment 1 (PHASE_SEG1)**: immediately precedes the sample point; can be lengthened by the resynchronization mechanism to absorb positive phase errors.
- **Phase Segment 2 (PHASE_SEG2)**: follows the sample point to the end of the bit; can be shortened to absorb negative phase errors.

The **sample point** falls at the PHASE_SEG1 / PHASE_SEG2 boundary. Every receiver samples the bus exactly once per bit at this point; the sampled polarity is the received bit value. The sample point position - expressed as a percentage of the total bit time - is a configuration parameter traded off against bus length, node count, and oscillator tolerance.

Resynchronization corrects for accumulated phase error between a receiver's local oscillator and the transmitter's bus edges. Hard synchronization forces a full re-alignment on the SOF falling edge at the start of each frame. During the frame, soft resynchronization adjusts PHASE_SEG1 or PHASE_SEG2 by up to the configured Synchronization Jump Width (SJW) on each recessive-to-dominant edge, keeping the sample point aligned with the transmitter. Only one synchronization is permitted within a single bit time (REQ-027). The two soft-resynchronization cases are illustrated in Figure 6.

CAN FD and the flexible data rate. CAN FD introduces a second, independently configured bit rate for the data phase. The BRS (Bit Rate Switch) bit in the FD control field (Figure 4) controls this transition: when BRS is recessive, the bus switches to the data-phase bit rate immediately after the BRS sample point and returns to the nominal rate at the CRC delimiter. The nominal rate governs the arbitration phase (SOF through BRS) and the return path (CRC delimiter onward); the data rate governs the payload and CRC fields in between. Because the data phase operates at a much shorter bit time, the same physical propagation delay represents a larger fraction of the bit period. On electrically long buses at high data rates, the loop propagation delay can exceed a full data-phase bit time.

Transmitter Delay Compensation (TDC) addresses this. A transmitter in the FD data phase cannot rely on immediate bus loopback for bit-error monitoring, because the echo of a driven bit arrives one or more bit times late. TDC measures the actual round-trip delay at the start of the data phase and configures a Secondary Sample Point (SSP) at the correct offset, so that each transmitted bit is still checked for loopback correctness. The TDC measurement and SSP configuration are PCS responsibilities and are a significant driver of PCS complexity in the implementation (Section ??).




Figure 4: Frame and error flag formats for the four structural frame formats (CB, CE, FB, FE) and the active/passive error flags. Remote frames (CB-RTR, CE-RTR) share the CB and CE structures respectively with an empty data field and RTR=1. Grey fields are fixed-polarity form fields; white fields carry data. Active and passive error flags each consist of six consecutive bits (dominant or recessive respectively), followed by an eight-bit recessive error delimiter ([REQ-007](#)). FD frames add BRS and ESI control bits and a longer payload and CRC sequence relative to Classic frames. Field widths follow ISO 11898-1 [\[1\]](#).




Figure 5: CAN bit time structure. One bit consists of SYNC_SEG (SS) which is 1 Time Quantum (TQ) long, PROP_SEG (PS), PHASE_SEG1 (PS1), and PHASE_SEG2 (PS2). The sample point (SP) sits at the PHASE_SEG1/PHASE_SEG2 boundary. The figure illustrates a to-node synchronised bus with bus wire delay t_{bus} and transceiver delay t_{TRX} .




Figure 6: Soft resynchronization over two successive sync edges. The phase error (PE) is the displacement between the received recessive-to-dominant edge and the start of SYNC_SEG. When PE exceeds SJW, PHASE_SEG1 is lengthened by exactly SJW quanta (After sync 1), moving the sample point forward by SJW. When PE is at most SJW, bit timing is restarted from SYNC_SEG on the incoming edge - equivalent to a hard synchronization - absorbing the full phase error in one step (After sync 2). After two such corrections the node is fully synchronised to the transmitter.




Figure 7: Transmitter Delay Compensation (TDC). At high data-phase bit rates the sum of bus propagation and transceiver delays (t_{loop}) may exceed one nominal bit time, causing the echo of a driven bit to arrive too late for the standard sample point. TDC measures this loop delay during the first data-phase bit and positions the Secondary Sample Point (SSP) at offset $t_{SSP} = t_{TDC_offset} + t_{measured}$, so that each transmitted bit is verified against its bus echo at the correct moment. The SSP replaces the nominal SP for bit-error monitoring throughout the data phase [1, Sec. 7.3.4].

4.4 Bit Stuffing

Bit stuffing ensures sufficient transitions on the bus for receiver clock synchronization. Classic CAN applies dynamic stuffing throughout the frame: after five consecutive bits of the same polarity, the transmitter inserts one complement stuff bit and the receiver removes it before forwarding the data stream (REQ-019). CAN FD retains dynamic stuffing through the arbitration phase, then switches to a combined dynamic-plus-fixed scheme in the data phase. Fixed stuff bits are inserted at predetermined positions (every fourth bit in the CRC field, independent of the preceding bit pattern); they carry a parity-encoded Stuff Bit Count (SBC) field that allows receivers to independently verify the number of dynamic stuff bits seen in the frame - an additional error detection layer absent in Classic CAN (REQ-017).




Figure 8: Dynamically and statically bit-stuffed stream examples. In frame fields encoded with dynamic bit-stuffing an opposite polarity stuff bit (SB) is inserted after five consecutive same-polarity bits. In frame fields encoded with static bit-stuffing (SBC and CRC in FD frames) a fixed stuff bit (FSB) is inserted after each fourth bit. A FSB is also inserted before the first bit of the SBC field.

4.5 Cyclic Redundancy Check

The CRC polynomial and field length depend on frame type and data payload length (REQ-006):

- **CRC-15:** used for all Classic CAN frames. Dynamic stuff bits are excluded from the CC CRC computation - the CRC accumulates over the de-stuffed bit stream from SOF through the end of the data field (REQ-013).
- **CRC-17:** used for FD frames with data payloads up to 16 bytes (DLC 0-10).
- **CRC-21:** used for FD frames with data payloads from 20 to 64 bytes (DLC 11-15).

A key asymmetry between CC and FD concerns the CRC data feed. For CC frames, dynamic stuff bits are excluded from the CRC computation (REQ-013). For FD frames, dynamic stuff bits within the arbitration region are included, while fixed stuff bits in the FD CRC region are also included. This dual data-feed requirement is a direct consequence of REQ-013 and has concrete consequences for the MAC implementation described in Section ???. The CRC field is terminated by a recessive CRC delimiter bit (REQ-018). A received frame with a CRC mismatch causes the detecting node to transmit an Error Flag.

4.6 Error Detection and Fault Confinement

Every CAN node monitors the bus for five categories of error (REQ-022). Bit errors occur when a transmitter reads back a polarity different from what it drove. Stuff errors occur when six consecutive bits of the same polarity appear where the stuffing rule prohibits it. CRC errors occur when the received checksum does not match the locally recomputed value. Form errors occur when fixed-format fields contain illegal bit values. Acknowledgement errors occur when a transmitter receives no dominant ACK bit from any receiver (REQ-014). Detection of any of these errors causes the detecting node to immediately transmit an Error Flag, aborting the in-progress frame (REQ-023). An Error Active node transmits an active error flag consisting of six consecutive dominant bits; an Error Passive node transmits a passive flag of six consecutive recessive bits instead (REQ-007). Both are followed by an eight-bit recessive error delimiter.

The FCE tracks each node's error history through TEC and REC. Counter increments and decrements follow the rules in ISO 11898-1 sec. 8.1.4.2 (REQ-030). A node begins in Error Active and transitions to Error Passive when either counter exceeds 127 (REQ-031), then to Bus Off when TEC exceeds 255. In Bus Off the node ceases all bus activity and shall not influence the bus (REQ-028) until 128 sequences of 11 consecutive recessive bits are observed, after which TEC and REC are reset and the node returns to Error Active. A host-initiated `llc_i.normal_mode` assertion also resets the FCE to its initial state immediately (REQ-029). Whether error signalling is enabled at all is a run-time configuration parameter (REQ-037). This escalation mechanism is the subject of several verification plan requirements and directly motivates the separation of the FCE into a dedicated module with its own testbench.

With those mechanisms established - sub-layer boundaries, frame formats, bit timing and the dual data rate, stuffing rules, CRC polynomials, and the fault confinement escalation model - the following section introduces the five classification dimensions of the verification plan and shows how each one connects back to the protocol concepts described here.

5 Verification Plan

The requirements set established what must be true about the implementation - 38 entries, each naming a protocol obligation and linking it to its ISO clause. But requirements in that form are not yet actionable as verification tasks: they say nothing about which module testbench should exercise them, what stimulus configurations are needed, whether internal signals must be observable, or how completion will be recognised. Turning the requirements set into a verification plan means answering those questions explicitly for each entry, before implementation begins.

The plan was populated through the same MCP server introduced in Section 3, which validated each field value against the schema before committing. The five classification dimensions fall into two groups. Three are design-facing - `layer`, `side`, and `format_applicability` - determining where each requirement belongs in the module decomposition and what stimulus configurations its testbench needs. Two are verification-facing - `observability` and `verification_method` - resolving whether a requirement can be checked through port signals or requires access to internal state, and specifying the verification technique. Priority spans both groups, driving implementation sequencing and determining which requirements must be closed before the design is considered complete. The dimensions are:

- `layer`, Section 5.1
- `side`, Section 5.2
- `format_applicability`, Section 5.3
- `observability`, Section 5.4
- `priority`, Section 5.5

The following subsections describe the rationale and allowed values for each dimension, followed by the full verification plan data structure and its traceability fields.

5.1 Layer

The `layer` field assigns each requirement to the protocol sub-layer that owns it (LLC, MAC, PCS, or FCE - see Section 4.1), determining the verification boundary at which the requirement must be exercised. A fifth label - **system** - classifies requirements that are inherently multi-layer or multi-node in character. Some CAN behaviors cannot be attributed to a single layer of a single node: they emerge from interactions between multiple nodes on the bus, or span the layer boundary within a single node. The `system` label

flags these requirements as ones that require either an integrated multi-module testbench or a multi-node simulation environment.

At design time, this classification directly motivated the layered module architecture: requirements assigned to a given layer pointed to the corresponding module as the responsible implementation unit and to that module's testbench as the primary verification environment. The consequence of the system label is described further in Section 6.2.2.

5.2 Side

The side field records whether a requirement pertains to the transmitter path, the receiver path, or both roles simultaneously. This dimension reflects the ISO standard's own framing, which frequently specifies transmitter and receiver obligations separately. In the verification environment, the side field determines whether a testbench drives the DUT in transmitter mode, receiver mode, or both roles in succession within a single test scenario. The design consequences of this dimension - and why it appeared to motivate a split-path architecture but did not - are discussed in Section 6.2.2.

5.3 Format Applicability

The format_applicability field records which of the in-scope frame formats (CB, CE, FB, FE - see Figure 4; CB and CE implicitly cover remote frame variants) each requirement applies to. Because the formats differ in stuffing mode, CRC polynomial, and control field structure (Section 4), a requirement that applies only to FD frames implies stimulus configurations with FDF=1 and DLC values spanning both the CRC-17 and CRC-21 threshold, while a requirement that applies to all four formats must be exercised across all format-specific configurations. The field makes those implications explicit rather than leaving them to be inferred from the requirement text.

5.4 Observability

The observability field resolves each requirement as either black-box or white-box, relative to the module boundary of the owning layer:

- **Black-box:** Can be verified purely through the module's observable port signals.
- **White-box:** Verification requires direct observation of the module's internal state.

This distinction has direct consequences for testbench architecture. Black-box requirements are verifiable with stimulus-and-observe testbenches that drive inputs and check outputs without any knowledge of internal implementation. White-box requirements - which include CRC polynomial correctness, bit counter arithmetic, error counter thresholds, and Gray-coded SBC encoding - require either PSL assertions on internal signals or a parallel reference model that re-computes the expected value independently. In the verification plan, white-box requirements are the primary driver for embedding PSL assertions directly in the RTL source files, where they have access to internal signals regardless of module hierarchy.

5.5 Priority

The priority field classifies each requirement into one of three levels:

- P1 requirements are need-to-have - they must be verified before the design can be considered complete.

- P2 requirements are nice-to-have - they are verified in the normal verification cycle but do not block closure.
- P3 requirements are optional - addressed only if schedule permits.

The final plan contains 31 P1, 5 P2, and 2 P3 requirements. The demotion rationale for each requirement not rated P1 is given in Table 3. The plan is considered closed when all P1 requirements reach complete status; P2 and P3 requirements are addressed as schedule permits.

5.6 Requirement Distribution

Table 4 shows the 38 requirements distributed across layer and observability. MAC dominates both in count and in white-box density, reflecting the breadth of frame-encoding logic that must be verified against internal bit-level state. FCE requirements are entirely black-box: fault-confinement state transitions are fully observable through the node’s error-state output signals without needing access to internal counters. System requirements - those that require a multi-node or multi-module environment - split roughly evenly between the two observability classes.

Table 4: Requirement distribution by layer and observability.

Layer	Black-box	White-box	Total
MAC	3	16	19
LLC	4	3	7
System	2	3	5
PCS	1	3	4
FCE	3	0	3
Total	13	25	38

5.7 Verification Plan Data Structure

The verification plan data structure (Table 5) augments each requirement with the dimensions needed to answer not just *what* must be true, but *how* it will be verified, *where* the evidence lives, and *when* verification is complete. The plan evolved continuously as implementation and verification work progressed. Two dimensions were not part of the initial taxonomy: the `system` layer label was added when it became clear that some CAN behaviors emerge from multi-node interactions and cannot be attributed to any single module’s testbench; the `observability` field was introduced when the distinction between black-box and white-box verification had direct consequences for testbench architecture that were not apparent from the requirement text alone. The following sub-sections detail the rationale and allowed values for each field in the verification plan. The complete verification plan is reproduced in Section A as two separate tables (linked by common ID’s).

Table 5: Verification-plan data structure fields.

Field	Purpose
<code>id</code>	Sequential identifier REQ-NNN.
<code>source_clause</code>	ISO 11898-1:2024 section reference.
<code>original_wording</code>	Verbatim normative text excerpts from the ISO standard.

Field	Purpose
<code>paraphrase</code>	Concise paraphrase of the <code>original_wording</code> field (this is the actual requirement)
<code>layer</code>	Sub-layer owner: LLC, MAC, PCS, FCE, or system (Section 5.1).
<code>side</code>	transmitter, receiver, or both (Section 5.2).
<code>format_applicability</code>	Applicable frame formats: CB, CE, FB, FE (Section 5.3).
<code>observability</code>	<code>black_box</code> or <code>white_box</code> (Section 5.4).
<code>verification_method</code>	Method(s) used to verify the requirement (Section 5.7.1).
<code>priority</code>	P1 (need-to-have), P2 (nice-to-have), or P3 (optional) (Section 5.5).
<code>status</code>	<code>not_started</code> , <code>in_progress</code> or <code>complete</code> (Section 5.7.3).
<code>notes</code>	The field is intended to clarify residual ambiguity left over from the other fields
<code>label</code>	Assertion label, TB procedure name, coverage ID, or RTL tag. Comma-separated when multiple procedures cover distinct sub-claims (Section 5.7.2).
<code>file</code>	Target file: TB for simulation/coverage, RTL for code inspection. Comma-separated when sub-claims span multiple files (Section 5.7.2).

5.7.1 Verification Method

The `verification_method` field makes the path from requirement to verification artifact explicit and actionable. Four methods are used: `simulation` (automated assertion procedures in a test bench), `code_inspection` (RTL source review), `waveform_inspection` (manual review of simulation output), and `coverage` (coverage bins a value range). Combinations are valid when multiple sub-claims within one requirement each call for a different method.

5.7.2 Traceability: Label and File

Each requirement entry carries two dedicated traceability fields: a `file` field identifying the test bench or RTL source file responsible for covering the requirement, and a `label` field identifying a specific named procedure, assertion, or coverage ID within that file. Together they establish a direct, navigable link from each requirement to its verification artifact.

5.7.3 Status

The `status` field (`not_started`, `in_progress`, `complete`) records requirement closure state explicitly, allowing partial progress to be tracked.

The verification plan - 38 requirements each classified along five dimensions and each linked to a testbench file and assertion label - served both roles in what followed. The design-facing dimensions (`layer`, `side`, `format_applicability`) constituted the primary architectural inputs for the design phase, mapping requirements to module boundaries and implementation scope. The verification-facing dimensions (`observability`, `verification_method`, `label`, `file`) defined the testbench architecture and evidence type for each requirement. How

the design-facing dimensions shaped the module decomposition - and where the apparent mapping from requirements structure to design structure broke down - is the subject of the next section.

6 Design and Architecture

This section traces how the three design-facing dimensions of the verification plan - `layer`, `side`, and `format_applicability` - shaped the module decomposition, and identifies where following the requirements structure led to a design decision that had to be revised. The final architecture and its module interfaces are then described in enough detail to motivate the implementation choices covered in Section 7.

6.1 Ramifications of the Requirements Model on Initial Design Strategy

The structure of the requirements model had direct consequences for the initial design strategy, in ways that were not fully anticipated at the outset.

The **layer dimension** mapped naturally onto the ISO standard's own layered reference model, making a layered module architecture look like the obvious and well-motivated implementation strategy. A dedicated hardware module for each layer - MAC, LLC, fault confinement, and PCS - would allow requirements pertaining to a given layer to be verified in isolation against that layer's module, rather than through the surface of a fully integrated system where internal behavior is obscured by surrounding logic. The observability dimension reinforced this directly: black-box requirements mapped cleanly onto port-level stimulus and observation, while white-box requirements pointed toward the need for reference models or PSL assertions on internal signals, both of which are most tractable in a per-module testbench. In retrospect, this was a sound conclusion. The modular architecture proved to be the right design choice, and the requirements table provided a well-motivated rationale for it from the start.

The **TX/RX side dimension** had a subtler and more consequential effect. Organizing requirements along the transmitter/receiver axis made intuitive sense from a specification perspective - the ISO standard itself frames many requirements in terms of transmitter behavior and receiver behavior - and it was genuinely useful for thinking through which requirements belonged where. However, it also made a split-path implementation architecture look like the natural design strategy, simply because the requirements were literally organized along that split. The implication appeared to be: implement a TX module, implement an RX module, and map the TX requirements to the former and the RX requirements to the latter.

This turned out to be a red herring. The pitfalls of the split-path approach were not at all apparent from the requirements table alone. The table made the split architecture look clean and well-motivated. The problems - design drift between separately implemented FSMs, integration complexity, and unnecessary hardware duplication - only surfaced later, during integration. This is an important general lesson: the structure of a requirements model can inadvertently bias architectural decisions in ways that are not immediately obvious, and the apparent naturalness of a design strategy that mirrors the requirements structure is not in itself a reliable signal that the strategy is sound.

The **format_applicability dimension** had a more constructive effect. Requirements tagged with a specific format subset (CB, CE, FB, FE, or combinations) made explicit which frame variants each protocol mechanism must handle, which in turn shaped two concrete design decisions. First, the MAC FSM needed per-field state granularity rather than per-phase granularity: because different format subsets diverge at specific field boundaries (the IDE/FDF/RES sequence differs between Classic and FD, and between base and extended addressing), encoding those branches in per-field states keeps each state's logic narrow and format-specific transitions visible in the state graph. Second, the LLC-to-MAC streaming interface needed

a front-loaded metadata layout - all format flags available before the first ID bit is needed - so that the MAC FSM could determine its branch path without buffering the entire frame. Both decisions are detailed in Section 6.2.3 and Section 6.2.4 respectively.

6.2 Architectural Design Decisions

The ramifications identified above narrowed the design space early: a layered architecture was well-motivated by the requirements model, and a unified FSM proved necessary once the split-path approach was attempted. The decisions documented in this section concern what remained open within that narrowed space. The primary inputs were the existing in-house CAN Classic controller (Section 1.2) and the ISO 11898-1 standard's own layered reference model [1]; CTU CAN FD [2], [3] is noted as an existing open-source CAN FD implementation but was not studied in detail.

6.2.1 Monolithic vs. Layered Architecture

The most fundamental architectural decision was whether to extend the existing monolithic controller or adopt a layered decomposition.

The existing controller uses a flat architecture: a single FSM (`can_fsm`, 810 lines) directly drives the bus output, reads the bus input, manages bit counting, handles stuff bit insertion, coordinates CRC computation, and updates error counters - all in one clocked process. This approach minimizes inter-module latency and signal fanout, since every decision is made in a single combinational cone. For CAN Classic, where the frame has at most two format variants (base and extended) and a single bit rate, the complexity is manageable.

CAN FD, however, introduces six bus frame variants (CB, CE, FB, FE data frames, plus remote frames for CB and CE), dual bit rates, fixed bit stuffing with SBC encoding, three CRC polynomials, and a richer error model. Extending the monolithic FSM to cover these features would require nested conditionals on the format type in nearly every state, increasing the cyclomatic complexity well beyond what is practical to verify or maintain. The existing controller's `s_arbitration` state already contains 70 lines of interleaved TX/RX logic for two frame formats; scaling this to six variants with additional control fields (BRS, ESI, SBC) and remote frame handling would roughly triple the state's complexity.

The alternative - and the approach adopted - is to follow the ISO 11898-1 reference model, which decomposes the data link layer into LLC, MAC, and PCS sub-layers with a separate FCE. Each sub-layer has a well-defined service interface (described in Section ??), and each can be implemented and verified independently. This decomposition introduces inter-module interfaces and pipeline latency, but it confines format-specific complexity to the module where it belongs: the MAC FSM handles frame field sequencing, the bit stuffer handles stuff bit insertion and SBC generation, and the CRC engine handles polynomial computation. No module needs to know about all three concerns simultaneously.

For the narrower scope of this project (protocol engine only, no message RAM or filtering), the full ISO layered decomposition is feasible and preferred because it directly maps each module to a testable subset of the standard's requirements.

6.2.2 Combined vs. Separated TX/RX FSMs

The initial design for the new MAC used **separate TX and RX FSMs**: `can_mac_fsm_tx` (~700 lines, 21 states) wrapped inside `can_mac_tx`, and `can_mac_fsm_rx` (~640 lines, 19 states) wrapped inside `can_mac_rx`. Each wrapper instantiated its own `can_mac_bs` and `can_mac_crc`. The two FSMs shared no state. The

only coupling was a `transmitting_i` flag passed from TX to RX, and a dominant-wins OR-merge of their respective PCS outputs at the `can_mac` wrapper level.

The argument for the split was grounded in the verification plan structure. The AI-assisted extraction described in Section 3.3 had classified each requirement by side (TX, RX, both), layer, and frame format. Mapping that classification onto separate entities looked elegant: TX-side requirements would be verified by exercising `can_mac_tx` in isolation, RX-side requirements by exercising `can_mac_rx` in isolation, with no cross-path stimulus needed. The separation also appeared to offer independence - a bug in one path could not corrupt the other's state.

In practice the split created more problems than it solved. Frame structure - the field ordering, the bit stuffing rules, and the CRC polynomials - is identical regardless of which node is driving. Encoding it twice meant that a fix to FD CRC delimiter handling in the TX FSM had to be replicated in the RX FSM with no compiler enforcement that the copies remained in sync. The doubled submodule footprint similarly doubled investigation surface: every "is the bit stuffer handling fixed stuffing correctly?" question had to be answered for two independent instances.

The most concrete cost was a debugging episode in `can_mac_pcs_fce_tb`, where a node that had successfully transmitted a frame was reporting `c_disturbed` instead of `c_transmitted`. Tracing the bug required correlating TX FSM state, TX bit count, RX FSM state, RX bit count, BS state on both sides, CRC state on both sides, and the merged PCS output - all in the same waveform pane, across two parallel state vectors that re-derived the frame position independently. The session made clear that single-bit-time bugs need a single-bit-time view, and that two parallel FSMs fundamentally opposed that.

The deeper lesson concerns the relationship between verification plan structure and RTL structure. Verification plan dimensions - TX/RX side, layer, frame format - are inputs to **testbench architecture**, not to RTL architecture. They describe what must be tested and what stimulus configuration is needed to test it. RTL architecture should follow the structure of the protocol itself. The natural code unit of the MAC is the frame, not the side, because a frame has the same shape whether the local node is driving it or sampling it. Cutting the natural code unit at an artificial seam, then re-joining the halves with shared submodules and a dominant-wins merge, added coupling without removing complexity.

The **final design** uses a single unified `can_mac_fsm` entity: one main FSM process (`p_fsm`), one `t_fsm_state` enum (19 states), one shared `can_mac_bs`, one shared `can_mac_crc`, and one `is_transmitter` mode flag latched at Start-of-Frame. TX-only requirements are verified with `is_transmitter = true` stimulus, RX-only with `is_transmitter = false` - the verification plan dimensions map cleanly onto testbench configurations rather than onto separate RTL entities.

6.2.3 Per-Field vs. Per-Phase FSM Granularity

A second FSM design decision was the granularity of states. The existing controller uses per-phase states: `s_arbitration` covers all ID bits, SRR, IDE, and RTR; `s_control` covers reserved bits and DLC; `s_data` covers all data bytes. Within each state, a bit counter and conditional logic dispatch the correct action for each bit position.

This per-phase approach keeps the state count low (18 states in the existing controller) but pushes complexity into the bit-counting logic. The `s_arbitration` state must distinguish between base and extended formats using the bit counter, handle the SRR/IDE branching point at bit 12/13, and detect arbitration loss - all in a single state with a single counter.

The alternative is per-field states, where each protocol field (SOF, ID, SRR/RRS, IDE, FDF, RES, BRS, ESI, DLC, Data, SBC, CRC, ACK, EOF) has its own state. This increases the state count - the unified FSM has 19 states - but eliminates most bit-counter conditionals: when the FSM is in `s_brs`, it knows exactly which bit it is processing without consulting a counter. Format-dependent transitions are handled by the state graph itself - for example, the FSM transitions from `s_ide` to `s_fdf_r1_r0` for all formats, but from `s_fdf_r1_r0` it may go to `s_dlc` (Classic) or `s_res_r0` then `s_brs` (FD). Each state contains only the logic relevant to its specific field, and named guard predicates (`v_in_arbitration_field`, `v_in_dynamic_stuff_field`, `v_in_fixed_stuff_field`) replace repeated bit-range comparisons.

The per-field approach was chosen because it aligns with the verification plan: each field in the frame structure corresponds to a set of requirements in the verification plan, and each FSM state can be traced directly to those requirements. It also simplifies formal verification, since PSL assertions can reference state names rather than counter ranges.

6.2.4 Internal LLC Frame Format

The user-facing LLC frame (shown in Figure 10) places all control flags at the end of the frame: FDF, BRS, and ESI occupy byte 69, and IDE and RTR occupy byte 70 - after up to 64 bytes of payload. A serializer that consumed the LLC frame in field order would need to buffer the entire 71-byte frame before it could begin transmitting, because IDE determines how many ID bits to drive (11 or 29), FDF determines which CRC polynomial and stuffing mode to use, and BRS determines whether to signal the PCS to switch bit rate at the BRS boundary. This buffering requirement conflicts with the streaming architecture.

The design avoids this by defining a separate internal format for the MAC-facing stream. All frame metadata is packed into two leading config bytes, followed by the ID and data bytes:

Byte	Bits	Field
0	[7]	IDE
0	[6]	FDF
0	[5]	FTYP (RTR)
0	[4]	ESI
0	[3]	BRS
1	[7:4]	DLC
1	[3:0]	(reserved, zero)
2-5	[7:0]	ID (32-bit, MSB first)
6+	[7:0]	Data (DLC bytes)

With this layout, `can_mac_ser` extracts all frame metadata after receiving just two bytes and can begin streaming ID bits from the third byte onward. No frame buffering is needed: the MAC FSM receives each metadata field before it is required in the frame field sequence, and the serializer operates as a pure pipeline stage.

6.3 System Overview

Figure 9 shows the complete module decomposition. The primary data path runs from `can_llc` through `can_mac` to `can_pcs`: the LLC receives frames from the host over an Avalon-ST interface and streams them byte-by-byte to the MAC serializer; the MAC FSM drives the serialized bit stream to the PCS, which

applies bit timing and produces the sample-point and SSP strobes that the MAC uses to read and write the bus. Two modules connect transversally across this pipeline: `can_fce` receives error and success events from the MAC and feeds node-state signals (error active, bus off) back to both the MAC and PCS; `can_pcs` also signals idle conditions to `can_fce` to support bus-off recovery. A centralized types package (`can_types_pkg`) defines all protocol constants, interface records, and reset values shared across modules.

6.4 LLC Frame Format

All six bus frame types (CB, CE, FB, FE data frames, and remote frames for CB and CE) are represented at the host-MAC boundary using the 71-byte LLC frame format shown in Figure 10. The LLC frame maps all in-scope variants into a fixed-width structure compatible with the Avalon-ST streaming interface, distinguishing remote frames via the FTYP bit in byte 70.

6.5 LLC Sub-layer

`can_llc` provides the Avalon-ST host interface and owns two protocol responsibilities above the MAC layer: retransmission management and acceptance filtering.

On the TX path, `can_llc` buffers one LLC frame from the host, streams it byte-by-byte to `can_mac_ser`, and monitors the transfer status returned by the MAC. On disturbance or lost arbitration, it retries up to the ISO-mandated limit before reporting failure to the host [1, Sec. 6.4.5] and 6.5. The MAC reports each outcome (transmitted, aborted, lost arbitration, disturbed); `can_llc` owns the retry counter and policy - the MAC is stateless with respect to retry.

On the RX path, `can_llc` receives completed LLC frames from `can_mac_fsm` over the Avalon-ST interface, applies acceptance filtering against the configured ID mask, and forwards accepted frames to the host [1, Sec. 6.4.5]. Frames that do not pass the filter are silently dropped.

The interface contracts `can_llc` must satisfy are captured in REQ-001 through REQ-005 and REQ-033; the module is not yet implemented.

6.6 MAC Sub-layer

The MAC sub-layer is the core of the protocol logic, responsible for bit serialization, CRC generation, bit stuffing, and frame-level error detection. It is implemented as a single unified `can_mac_fsm` entity, supported by three internal sub-modules (`can_mac_ser`, `can_mac_bs`, `can_mac_crc`) and wrapped by `can_mac`, which is a structural entity that instantiates the FSM alongside `can_fce` and exposes their combined LLC, PCS, and FCE interfaces. It coordinates closely with the FCE (Section 6.7) for error counter management and node-state transitions (Error Active/Passive/Bus Off), and with the PCS (Section 6.8) for sample-point-driven bit output.

6.6.1 `can_mac_fsm`

`can_mac_fsm` is a 19-state per-field FSM that handles both frame transmission and frame reception. An `is_transmitter` flag latched at SOF partitions per-state logic into TX and RX branches without duplicating the state graph. In TX mode, the FSM drives bits through the PCS at each bit-boundary strobe, monitors the bus echo for bit errors and arbitration loss, and feeds the bit stuffer and CRC engine. In RX mode, it observes sampled bus bits, performs destuffing, accumulates the CRC, validates form fields, drives the ACK slot, and streams the completed frame to the LLC during intermission. Errors in either mode branch




Figure 9: Implementation module decomposition. `can_llc` provides the Avalon-ST host interface and manages retransmission and acceptance filtering. Inside `can_mac`, the unified `can_mac_fsm` handles both TX and RX roles via a single `is_transmitter` flag, sharing one `can_mac_bs` bit stuffer and one `can_mac_crc` engine. `can_pcs` generates sample-point and SSP strobes. `can_fce` maintains TEC/REC counters and governs node-state transitions.




Figure 10: LLC frame format (71 bytes) at the host-MAC interface. Bytes 0-3 encode the identifier; byte 4 carries the frame-type selector FMT in bits [7:4] and DLC in bits [3:0]; bytes 5-68 carry up to 64 data bytes; byte 69 carries the FDF, BRS, ESI, and IDE control flags; byte 70 carries RTR in bit 0. The lower rows show the identifier byte mapping: for extended identifiers ID[28:24] are packed into ID0, ID[23:16] into ID1, ID[15:8] into ID2, and ID[7:0] into ID3; for base identifiers ID0 and ID1 are all zeros, ID[10:8] occupies the lower three bits of ID2, and ID[7:0] fills ID3.

to a two-state error-frame sequence (`s_error_flag`, `s_error_delimiter`). Detailed implementation is described in Section 7.2.

6.6.2 `can_mac_ser`

`can_mac_ser` converts the LLC byte stream into a serial polarity bit stream for the MAC FSM. It manages the two-byte configuration handshake (see Section 6.2.4), extracts LLC metadata (IDE, FDF, DLC, FTYP, BRS, ESI) from the config bytes, and serializes ID and data bits one per FSM ready pulse. Unused padding bits in the 32-bit ID field for 11-bit base identifiers are skipped silently. Detailed implementation is described in Section 7.3.

6.6.3 `can_mac_bs`

`can_mac_bs` implements both dynamic and fixed bit stuffing for CAN Classic and CAN-FD frames [1, Sec. 10.6], used for both TX stuffing and RX destuffing via the FSM's `is_transmitter` mode. In dynamic mode, it inserts an inverse-polarity bit after five consecutive identical bits. In fixed mode (FD CRC region), it inserts one bit on entry then one every four bits, and maintains a Gray-coded stuff bit count with parity for the SBC field. Detailed implementation is described in Section 7.4.

6.6.4 `can_mac_crc`

`can_mac_crc` runs three parallel `gen_crc` instances - CRC-15, CRC-17, and CRC-21 - on two independent data feeds: `data_cc` (de-stuffed, for Classic frames) and `data_fd` (includes dynamic stuff bits in the arbitration region, for FD frames). The active engine is selected by `crc_poly_select` and its result is left-aligned to a common 21-bit output. Detailed implementation is described in Section 7.5.

6.7 FCE Sub-layer

`can_fce` implements the error state machine and counter management specified in ISO 11898-1 [1, Secs. 8.1.3–8.1.4]. It maintains TEC and REC and transitions between three states: `s_error_active` (initial), `s_error_passive` (TEC or REC > 127), and `s_bus_off` (TEC > 255). Counter increment and decrement rules follow ISO 8.1.4.2. Bus-off recovery counts 128 `pcs_i.idle_condition` pulses from the PCS (11 con-

secutive recessive bits each), or responds to `llc_i.normal_mode` from the LLC. Detailed implementation is described in Section 7.6.

6.8 PCS Sub-layer

`can_pcs` handles bit timing, hard and soft synchronization, and Transmitter Delay Compensation (TDC) for both TX and RX paths [1, Secs. 7.2–7.3]. It is a cyclic bit-timing engine: a `t_segment` register cycles through `s_sync_seg`, `s_prop_seg`, `s_phase_seg1`, and `s_phase_seg2` each bit time. The SP strobe fires at the end of `s_phase_seg1`. Hard synchronization restarts the cycle on a dominant edge in `s_sync_seg`; soft synchronization adjusts `s_phase_seg1` or `s_phase_seg2` by up to SJW. At the BRS sample point the PCS switches to the independently configured data-phase segment lengths. TDC measures the TX-to-RX echo delay at the first data-phase bit and positions the SSP accordingly. When bus-off is asserted, the PCS counts consecutive recessive bits and pulses `fce_o.idle_condition` every 11 bits for FCE recovery. Detailed implementation is described in Section 7.7.

The decomposition described above yields five implemented entities - `can_mac_fsm`, `can_mac_ser`, `can_mac_bs`, `can_mac_crc`, `can_fce`, and `can_pcs` - plus `can_mac` and `can_mac_pcs_fce` as structural wrappers and `can_llc` as the one module not yet implemented. The following section covers the implementation of each entity in turn, ordered MAC-first.

7 Implementation

This section describes the implementation of the six modules that together constitute the CAN-FD controller: `can_mac_fsm`, `can_mac_ser`, `can_mac_bs`, `can_mac_crc`, `can_fce`, and `can_pcs`. `can_llc` is not yet implemented; its interface contracts are captured in the verification plan (REQ-001 through REQ-005, REQ-033). Subsections are ordered MAC-first: the central FSM and its internal submodules are described before the supporting FCE and PCS layers.

7.1 Interface Conventions

The company's `std_logic`-only port constraint does not preclude the use of VHDL record types on entity ports - records of `std_logic` and `std_logic_vector` fields satisfy the constraint. The design uses typed record interfaces (e.g., `t_can_mac_pcs_if_m2s`, `t_can_mac_fsm_bs_if_s2m`) to bundle related signals into a single port. Each record type has a corresponding reset constant (e.g., `c_can_mac_pcs_if_m2s_reset`), ensuring that every module can be reset to a known state without manually enumerating each field.

The alternative - flat port lists with individual `std_logic` signals - was used in the existing controller, where the FSM entity has 20 individual ports for its various sub-module interfaces. This approach becomes unwieldy as the number of inter-module signals grows: the new design has over 60 inter-module signals across its interfaces, and bundling them into records reduces the port list from dozens of lines to six record-typed ports per FSM entity.

The naming convention follows the data-flow direction: `m2s/s2m` (master-to-slave/slave-to-master) for control interfaces, and `s2d/d2s` (source-to-destination/destination-to-source) for Avalon-ST data-transfer interfaces. This convention is consistent with the company's existing interface naming and makes the direction of data flow explicit at every port.

7.2 can_mac_fsm

The `can_mac` sub-layer is built around a single unified FSM entity (`can_mac_fsm`, ~1100 lines). The full per-signal interface is shown in Figure 11.




Figure 11: ‘`can_mac`’ architecture. Individual signals are shown on each directed edge. MAC-internal connections: ‘`can_mac_ser`’ feeds the serial bit stream and ‘`llc_metadata`’ to ‘`can_mac_fsm`’; ‘`can_mac_bs`’ accepts ‘`data/valid/fsb_en`’ and returns stuffed ‘`data/valid`’ plus the Gray-coded ‘`stuff_bit_count`’ for the SBC field; ‘`can_mac_crc`’ receives the dual CC/FD bit feeds and returns the running ‘`crc[20:0]`’ digest. External connections: ‘`can_llc`’ connects to the MAC on both TX (‘`tx_llc_i/o`’ via the serializer) and RX (‘`rx_llc_i/o`’ via the FSM); bit-level timing to ‘`can_pcs`’ (‘`tx_data`’, ‘`rx_data`’, sample-point strobes); fault-confinement events to/from ‘`can_fce`’ (‘`error_active`’, ‘`bus_off`’).

7.2.1 FSM Structure and Mode Flag

`can_mac_fsm` contains two synchronous processes (`p_fsm` and `p_stream_to_LLC`) and one `t_fsm_state` enum covering 19 states. An `is_transmitter` boolean signal is latched to `true` when the FSM drives the SOF dominant bit at the start of a new frame transmission and cleared at arbitration loss or at the end of the EOF field. Once latched, `is_transmitter` remains stable for the rest of the frame, partitioning per-state logic into a TX branch (`drive_bit` strobe: determines and drives `pcs_o.tx_data`) and an RX branch (`sample_point` strobe: advances state and captures received bits). The state transitions at the sample point are shared between TX and RX in most states; only `s_ack`, `s_ack_delimiter`, and `s_eof` carry role-specific logic on the sample-point path (ACK success latching, receiver dominant assertion, and frame-completion signalling respectively).

The per-field state granularity introduced in Section 6.2.3 is preserved in the unified FSM. The arbitration region uses two states: SOF is driven in `s_bus_idle`, and the ID bits, RTR/SRR/RRS, and IDE share `s_arbitration` with `bit_count` indexing into the 32-bit ID field. Each post-arbitration protocol field (FDF, RES, BRS, ESI, DLC, Data, SBC, CRC, CRC Delimiter, ACK, ACK Delimiter, EOF) has a dedicated state, making field boundaries explicit in the state encoding. The complete FSM is shown in Figure 12.

7.2.2 TX Mode: Frame Transmission

When `is_transmitter = true`, the FSM drives bits through `pcs_o.tx_data` at the bit-boundary strobe (`drive_bit`, generated internally by registering `pcs_i.sample_point` twice). On each sample-point strobe the FSM executes the following four-step sequence per frame-field state:

1. **Monitor the bus.** The sampled bus polarity (`pcs_i.rx_data`) is compared against the previously transmitted bit (held in a 32-bit polarity history shift register, `transmitted_bits_shift_reg`) to detect bit errors, ACK, and arbitration loss. In the CAN-FD data phase, any pending SSP observation from the previous bit period is evaluated against this history before the current bit is checked [1, Sec. 7.3.4]. A detected error triggers a transition to the error-frame sequence with the appropriate flag type based on FCE fault confinement status [1, Secs. 8.1.3–8.1.4]. Arbitration loss causes `is_transmitter` to flip to false in-place, and the `s_arbitration` state then continues as an RX observer for the remaining bits.
2. **Determine the next bit.** If the bit stuffer has a pending stuff bit (`bs_i.valid`), that takes priority. Otherwise, the polarity is determined by the current state: form bits (SOF, IDE, FDF, reserved, delimiters, EOF) have fixed polarities, while ID, DLC, data, SBC, and CRC bits are sourced from the serializer, metadata, stuff bit count, or CRC register respectively.
3. **Feed the CRC engine and bit stuffer.** The post-case feed source is `pcs_i.rx_data` in `s_arbitration` (where the bus is multi-master and the winning transmitter's echo matches its drive on the wire) and `transmitted_bits_shift_reg(0)` in all subsequent states - both nominal and data phase. Using the shift register post-arbitration ensures the feed is independent of any `rx_data` echo latency under TDC delay and keeps the source consistent across all non-arbitration states. The FSM asserts `bs_o.fixed_bit_stuffing_en` when entering the SBC field of FD frames to switch the bit stuffer from dynamic to fixed mode.
4. **Present the bit at the PCS interface.** The resolved polarity is written to `pcs_o.tx_data` and becomes visible at `tx_o` when the bit-boundary latch fires. `pcs_o.next_bit_is_brs` is asserted one SP before the BRS bit, allowing the PCS to switch to data-phase timing at the BRS SP if BRS is sampled recessive. `pcs_o.next_bit_is_res` is asserted one SP before the FD reserved bit to arm TDC measurement at the subsequent bit boundary [1, Sec. 7.3.4].

7.2.3 RX Mode: Frame Reception

When `is_transmitter = false`, the FSM observes `pcs_i.rx_data` at each sample-point strobe and stores received bits directly into an internal `llc_frame` byte array. No separate deserializer is needed. The bit stuffer is driven from `pcs_i.rx_data` to perform destuffing, and the CRC engine accumulates the received bit stream in parallel. The FSM validates the SBC field (FD frames), compares the received CRC against the locally accumulated result, and checks form bits (reserved bits, CRC delimiter, ACK delimiter, EOF) for required polarities. A mismatch in any of these fields triggers a transition to the error-frame sequence. During the ACK slot the FSM drives `pcs_o.tx_data = c_dominant` for one bit (`bit_count = 0`) regardless of frame format; the FD ACK slot spans two bits but the receiver asserts dominant only during the first. The bus is released after the ACK delimiter [1, Sec. 8.1.4.2.b].

After the EOF field the FSM transitions through `s_intermission`. A dedicated second process, `p_stream_to_LLC`, transfers the completed `llc_frame` array byte-by-byte to the LLC RX sink over the Avalon-ST interface; `p_fsm` triggers the transfer by asserting `llc_stream_start` and separately signals successful reception to the FCE. This design eliminates the need for a separate deserializer entity on the

RX path - the frame buffer is populated and streamed entirely within `can_mac_fsm`.

7.2.4 Error-Frame States

The FSM uses two explicit error-frame states. `s_error_flag` drives the 6-bit flag and `s_error_delimiter` counts the 8-bit recessive delimiter. The delimiter state manages an internal phase flag (`delim_found_first_recessive`) that separates two distinct sub-phases: first awaiting the bus to go recessive (other nodes may still be driving their own flags), then counting the remaining recessive bits. A dominant during the delimiter restarts the error-frame sequence - either as a new error or as an overload condition on the last delimiter bit (ISO 8.1.4.2.f).

Table 6: Two explicit error-frame states of `can_mac_fsm`.

State	ISO reference	Role
<code>s_error_flag</code>	§6.6.5.2, §6.6.13	Drive 6 dominant bits (active/overload) or 6 recessive bits (passive).
<code>s_error_delimiter</code>	§6.6.5.3, §8.1.4.2	Await first recessive bit, then count 8-bit delimiter. A dominant restarts the error-frame sequence.

7.3 `can_mac_ser`

`can_mac_ser` converts the LLC byte stream into a serial polarity bit stream for the MAC FSM. Its four-state FSM manages the two-byte configuration handshake, byte fetching, and bit-by-bit serialization. The serializer extracts LLC metadata (IDE, FDF, DLC, FTYP, BRS, ESI) from the two config bytes and registers it in `t_llc_metadata`, which remains stable for the entire frame. The internal frame format that makes this possible is described in Section 6.2.4. The serializer forwards `transfer_status` from the FSM back to the LLC, returning to `s_load_config_byte_0` on any non-ongoing status so that errors and aborts terminate serialization immediately.

The 32-bit ID field in the internal format is right-aligned: an extended identifier (29 bits) uses bits [28:0], leaving 3 unused bits; a base identifier (11 bits) uses bits [10:0], leaving 21 unused bits. The serializer tracks this with two counters initialized in `s_load_config_byte_1` from the `ide` flag: `id_bits_remaining` counts real ID bits still to be presented, and `padding_bits_remaining` counts leading zeros to be skipped. In `s_shift_out_bits`, padding bits are advanced without asserting `valid`, so the MAC FSM never observes them. This means the FSM always receives exactly 11 or 29 consecutive valid ID bits regardless of frame format, with no format-specific logic required downstream.

Bit serialization uses a single-byte shift register. On entry to `s_shift_out_bits`, `llc_frame_buffer` holds the current byte with the MSB pre-loaded into `tx_mac_fsm_o.data`. The serializer holds `valid` high while a real bit is waiting; the FSM acknowledges by asserting `ready` for one cycle. On each accepted transfer the buffer is shifted left by one and the new MSB is presented as the next bit. When the final bit of the byte is consumed (`count = c_byte_width - 1`), `valid` is deasserted and the FSM returns to `s_load_llc_frame_byte` to fetch the next byte. The LLC's Avalon-ST `ready` signal is held low during `s_shift_out_bits`, preventing the LLC from advancing while a byte is being drained.




Figure 13: `can_mac_ser` (4 states). Config bytes carry frame metadata (IDE, FDF, FTYP, BRS, ESI, DLC); these are latched once and held stable for the entire frame. In `s_shift_out_bits` real ID and data bits are driven one per MAC FSM ready pulse, while unused padding bits in the 32-bit ID field are skipped silently. Any `transfer_status != ongoing` signal from the MAC FSM returns the serializer to `s_load_config_byte_0` from any state.

7.4 `can_mac_bs`

`can_mac_bs` implements both dynamic and fixed bit stuffing for CAN Classic and CAN-FD frames [1, Sec. 10.6]. The single entity is instantiated once inside `can_mac_fsm` and serves both TX stuffing and RX destuffing via the same logic - the FSM drives the same `bs_i` interface regardless of role, and the stuffer's output is either inserted into the TX bit stream or used by the FSM to discard a received destuff bit.

In **dynamic mode** (`fixed_bit_stuffing_en = '0'`), the stuffer counts consecutive bits of identical polarity and emits an inverse-polarity stuff bit after every five (REQ-019). The Gray-coded stuff bit count `stuff_count` is incremented on each dynamic stuff bit and encoded with parity into the `stuff_bit_count` output, which the FSM reads when transmitting the SBC field (REQ-017).

In **fixed mode** (`fixed_bit_stuffing_en = '1'`), used for the FD CRC region, one fixed stuff bit (FSB) is emitted immediately on the rising edge of `fixed_bit_stuffing_en`, then one FSB every four real bits (REQ-019). The FSB polarity is always the inverse of the preceding bit, so a receiver can detect a form error if the FSB matches its predecessor.

The transition from dynamic to fixed stuffing requires special handling when a dynamic stuff bit is already pending at the rising edge of `fixed_bit_stuffing_en`. Suppressing the pending dynamic SB would cause a TX/RX divergence: the transmitter and receiver would derive different `stuff_count` values from the same bit stream, causing SBC mismatch. The implementation instead promotes the pending dynamic SB to the initial FSB - the two coincide and both requirements are satisfied simultaneously (see Figure 14, REQ-017, REQ-019). On the falling edge of `fixed_bit_stuffing_en`, any pending FSB is cancelled immediately; the MAC FSM exits fixed stuffing at the last CRC bit without providing a slot to drain a still-pending

FSB.




Figure 14: can_mac_bs (2 modes). Dynamic mode counts consecutive identical bits and emits an inverse stuff bit on the 5th (count = 4), updating the Gray-coded SBC. Fixed mode is entered on the rising edge of fsb_en, which immediately emits the initial FSB; thereafter one FSB is emitted every 4 real bits (count = 3). The falling edge of fsb_en cancels any pending FSB and returns to dynamic mode. stuff_count and the SBC output are maintained continuously across both modes.

7.5 can_mac_crc

CAN Classic computes its CRC over the raw bit stream excluding stuff bits, while CAN FD includes dynamic stuff bits in the arbitration region - a difference that exists because the FD CRC must protect the stuff bit count as well as the data. A single data feed to the CRC engine is therefore insufficient: CC and FD frames require different input streams. The design exposes two feeds on the CRC interface (data_cc and data_fd) so the FSM can drive both simultaneously and the CRC module requires no protocol knowledge about which stream to select.

can_mac_crc provides CRC generation and checking for both CAN Classic and CAN-FD frames. CAN Classic frames use CRC-15, while CAN-FD frames use CRC-17 (data payloads up to 16 bytes) or CRC-21 (data payloads above 16 bytes) [1, Sec. 10.4.2.6]. The single entity is instantiated once inside can_mac_fsm, serving both TX (generation) and RX (checking). The FSM sets crc_poly_select from the DLC field in llc_metadata before the first frame bit is driven: because the internal LLC frame format (Section 6.2.4) delivers DLC in config byte 1, the polynomial is known upfront and requires no mid-frame switching.

Three parallel gen_crc instances run continuously on separate data feeds: data_cc drives CRC-15 via valid_cc, while data_fd drives both CRC-17 and CRC-21 via valid_fd. This dual-feed architecture is necessary because CC and FD frames compute CRC over different bit streams (CC excludes stuff bits; FD includes them in the arbitration region), and the RX path does not know which CRC engine to use until after the frame type has been determined. The output multiplexer selects the active engine's result based on crc_poly_select and left-aligns it to the common 21-bit output width by zero-extending the shorter results at the LSB: CRC-15 occupies bits [20:6], CRC-17 occupies bits [20:4], and CRC-21 occupies the full width.

The output mux (p_crc_mux) is implemented combinatorially rather than as a registered stage. Each gen_crc instance registers its accumulator on the rising edge; a registered mux would add one cycle of latency, causing the FSM to read a stale digest on the cycle it drives the first CRC bit. The combinatorial mux ensures crc_o.crc reflects the fully accumulated result on the same cycle that the last data bit is

registered, so the FSM can begin driving CRC bits immediately without an explicit wait state.




Figure 15: `can_mac_crc` dataflow. Three `gen_crc` instances accumulate in parallel on every clock: `u_crc15` is fed by `data_cc/valid_cc` (CC bit stream, stuff bits excluded); `u_crc17` and `u_crc21` are both fed by `data_fd/valid_fd` (FD bit stream, dynamic stuff bits included in the arbitration region). The combinatorial mux `p_crc_mux` selects the active engine result based on `crc_poly_select` and zero-extends it to the common 21-bit output. `reset_i` reinitializes all three engines simultaneously.

7.6 `can_fce`

`can_fce` implements the error state machine and counter management specified in ISO 11898-1 [1, Secs. 8.1.3–8.1.4]. It maintains TEC (Transmitter Error Counter) and REC (Receiver Error Counter) and transitions between three states: `s_error_active` (normal operation), `s_error_passive` (TEC or REC > 127), and `s_bus_off` (TEC > 255).

Counter updates follow the ISO 8.1.4.2 rules: TEC increments by 8 on TX errors, with `mac_i.passive_tx_ack_error_ex` suppressing the increment for the passive ACK error exemption (ISO 8.1.4.2.c, Exception 1); TEC decrements by 1 on successful TX. REC increments by 1 on RX errors during non-error-flag phases, by 8 on primary errors or error-flag-phase errors, and decrements by 1 or clamps to 127 on successful RX. Bus-off recovery requires counting 128 `pcs_i.idle_condition` pulses (11 consecutive recessive bits each) from the PCS, or assertion of `llc_i.normal_mode` by the LLC, either of which resets both counters and returns the FSM to `s_error_active`.




Figure 16: can_fce FSM (ISO 11898-1 sec. 8.1.4.4, Figure 43). `s_error_active` is the initial state after `reset_i` or `llc_i.normal_mode`. When TEC exceeds 127 or REC exceeds 127 the FCE moves to `s_error_passive`, clearing `mac_o.error_active`. If TEC exceeds 255 from either active or passive state, the FCE enters `s_bus_off`, asserting `bus_off` to the MAC, LLC, and PCS. Recovery counts 128 `idle_condition` pulses from the PCS (each pulse represents 11 consecutive recessive bits); on the 128th pulse TEC and REC are reset to zero and the FSM returns to `s_error_active`.

7.7 can_pcs

`can_pcs` is a cyclic bit-timing engine: its internal `t_segment` register advances through `s_sync_seg` (1 TQ, fixed), `s_prop_seg`, `s_phase_seg1`, and `s_phase_seg2` on every TQ boundary. The SP strobe and `rx_data` latch fire at the end of `s_phase_seg1`; the TX bit is driven at the end of `s_phase_seg2`. When `fce_i.bus_off` is asserted, the SP slot counts consecutive recessive bits and pulses `fce_o.idle_condition` every 11 bits for FCE bus-off recovery. The full timing operation is shown in Figure 17.

7.7.1 Resynchronisation

The prior implementation (`can_node_clock`) did not satisfy three of the four ISO 7.3.5.1 synchronisation rules: it had no sync-inhibit guard (rule a), no sampled-polarity check (rule b), and it contained a Sync_Seg skip defect in which Phase_Seg2 shortening caused the module to jump directly from `s_seg_2` to `s_seg_1`, dropping the mandatory 1-TQ Sync_Seg entirely. `can_pcs` corrects all three and enforces all four rules.

Rule a - one synchronisation per bit time. A `sync_applied` signal is set on any synchronisation event (hard or soft) and cleared at the next bit boundary (end of `s_phase_seg2`). The TQ-boundary edge-qualify predicate `v_do_sync` includes `sync_applied = '0'` as a precondition, preventing a second synchronisation within the same bit time regardless of bus activity.

Rule b - sync only on a recessive-to-dominant transition. The edge-qualify predicate requires `rx_bus_prev = c_recessive` (the bus value latched at the preceding TQ boundary) together with `rx_i = c_dominant`, making synchronisation conditional on an actual recessive-to-dominant edge. This prevents spurious synchronisation on a dominant-to-recessive-to-dominant glitch within a dominant bit. A further guard, `mac_i.transmitting = '0'`, disables synchronisation entirely while the local node is driving the bus.

Rule c - hard synchronisation on demand. Rather than re-entering a hard-sync state via a full module reset (as `can_node_clock` did via its `reset_i` port), `can_pcs` accepts a MAC-driven `mac_i.do_hard_sync` signal. When asserted, any qualifying edge triggers a full bit-time restart from `s_prop_seg` with the prescaler and segment counter cleared. This allows the MAC to switch synchronisation mode at any point - including at the FDF-to-res transition required by ISO 7.3.5.1(c) - without resetting PCS timing state.

Rule d - Sync_Seg always traversed. `s_sync_seg` is an unconditional stop in the segment FSM: `s_phase_seg2` always transitions to `s_sync_seg` at the bit boundary, and `s_sync_seg` always transitions to `s_prop_seg` after 1 TQ. No shortcut paths exist.

7.7.2 Dual Bit Rate Switching

`can_pcs` holds no frame-format knowledge. Rate switching is entirely MAC-driven through three dedicated control signals on the MAC-PCS interface.

`mac_i.next_bit_is_brs` is asserted one SP before the BRS bit. At the BRS SP, `can_pcs` reads `rx_i` (or `mac_i.tx_data` when transmitting) to determine BRS polarity; if recessive, it replaces `active_prop_seg`, `active_phase_seg1`, `active_phase_seg2`, and `active_sjw` with the data-phase generics (`gc_data_prop_seg`, `gc_data_phase_seg1`, `gc_data_phase_seg2`, `gc_data_sjw`) and sets `data_phase_active`. `mac_i.next_bit_is_res` is asserted when the next bit is the FD reserved bit; at the corresponding bit boundary, `tdc_count_active` is set to begin TDC measurement (described in Section 7.7.3). `mac_i.data_phase_stop` is asserted by the MAC at the CRC delimiter SP or on entry to the error-frame sequence; it restores nominal timing parameters and clears all TDC state.

This interface design keeps protocol knowledge in the MAC layer and timing knowledge in the PCS layer, following the ISO 11898-1 layered architecture. `can_pcs` does not inspect DLC or frame format - the only frame-level observation it makes is reading the BRS bit polarity when `mac_i.next_bit_is_brs` is set, making it independently testable against any timing configuration without frame-level stimulus.

7.7.3 Transmitter Delay Compensation

The motivation and principle of TDC are described in Section 4.3. The implementation in `can_pcs` is flag-based logic within the single `p_can_pcs` process, not a separate state machine. The relevant signals are `tdc_count_active`, `delay_count_tq`, `ssp_standoff_active`, `first_data_bit_boundary_seen`, `ssp_active`, `ssp_seen`, and `tdc_delay`.

TDC is armed at the bit boundary when `mac_i.next_bit_is_res = '1'`, which sets `tdc_count_active`. While set, `delay_count_tq` increments by one per recessive TQ at each TQ boundary. On the first dominant TQ - the TX-to-RX echo of the data-phase preamble - `tdc_count_active` is cleared, leaving `delay_count_tq` holding the measured round-trip delay in TQs.

At the first data-phase bit boundary (`data_phase_active = '1'`, `first_data_bit_boundary_seen = '0'`, `ssp_seen = '0'`), `first_data_bit_boundary_seen` is latched and `ssp_standoff_active` is asserted. From that point `delay_count_tq` counts down one per TQ; when it reaches zero, `ssp_active` is set and `ssp_standoff_active` is cleared. With `ssp_active = '1'`, the SSP fires one TQ before the SP in each subsequent bit time (at `seg_count = active_phase_seg1 + phase1_extension - 2` within `s_phase_seg1`), latching `rx_i` and pulsing `mac_o.secondary_sample_point`. On the first SSP, `ssp_seen` is latched to prevent repeated SSP firing within the same data phase. At each data-phase bit boundary after the first (while `ssp_seen = '0'`), `tdc_delay` increments to count whole bit-period delays; the MAC reads `mac_o.tdc_delay` to index into `transmitted_bits_shift_reg` and identify the transmitted bit corresponding to the received bit observed at the SSP.

`mac_i.data_phase_stop` at the SP clears `ssp_active`, `ssp_seen`, `tdc_count_active`, `delay_count_tq`, `tdc_delay`, `data_phase_active`, and `first_data_bit_boundary_seen`, restoring nominal SP-based monitoring for the CRC delimiter and subsequent fields.

8 Verification and Results




Figure 18: [REQ-009](#). Two-node simulation of a complete FD frame. DUT 1 is the transmitter (`is_transmitter = true`), DUT 2 the receiver. `mac_fsm_state` traces show the full frame-field sequence from `s_arbitration` through `s_eof`; `secondary_sample_point` pulses on DUT 1 confirm TDC is active during the data phase; `sync_applied` pulses on DUT 2 show receiver soft-synchronisation events. After EOF both nodes traverse `s_intermission` and the sample-point counter accumulates 11 consecutive recessive bits, producing the idle condition and completing bus re-integration ([REQ-009](#)).




Figure 19: [REQ-010/011](#). Error escalation waveform for DUT 1. `transmitter_error_count` (TEC) increments on successive error frames (13 → 128 → 136); when TEC exceeds 127, `error_active` deasserts and the node becomes error-passive. The `mac_fsm_state` trace shows `s_suspend_transmission` inserted after `s_intermission` once the node is error-passive, enforcing the mandatory 8-bit suspension period before the next transmission attempt (ISO 11898-1 §8.1.4.4).

Note: This section is currently being populated as the verification plan is executed.

8.1 Testbench Results Summary

Table 7: Testbench execution status and functional coverage summary.

Testbench	Status	Coverage
tx_pcs_tb	Pass	95%
tx_can_tb	Pass	88%

Testbench	Status	Coverage
tx_can_protocol_tb	Pass	92%

9 Discussion

Comparison of the implemented architecture against theoretical models. Performance analysis in high-load scenarios.

9.1 Future Work

1. Make the the CRC and BS modules to save area on the FPGA.
2. CAN XL implementation...

10 Conclusion

Summary of work completed and how objectives were met.

11 References

A Verification Plan

Both tables are regenerated automatically from `verification_plan/verification_plan.toml` on each PDF build. The ID field is the join key between them. See Section 5.7 for the meaning of each field.

A.1 Extracted Requirements

The requirements table contains four fields: the identifier, the ISO 11898-1 source clause, the priority assignment, and the paraphrase. Placing priority here allows each assignment to be evaluated directly against the requirement text. The table corresponds to the requirements extraction and prioritisation phases described in Section 3.2 and Section 5.5.

Table 8: ISO 11898-1 extracted requirements.

ID	ISO ref	Priority	Requirement	Notes
REQ-001	6.4.5.2	P1	LLC notification LLC shall notify the user on every successful DF/RF TX or RX.	Timestamping (§6.5.7) is optional per ISO and is out of scope.
REQ-002	6.4.5.5.2, 6.4.5.5.3	P2	LLC transmission request and abort timing 1) A TX request shall be processed no later than the second SOF after the request when no EFs are present. 2) Frames not yet passed to the MAC are aborted immediately. 3) An abort request shall be processed before the second SOF (no EFs present) or third SOF (with EFs present).	Frames already passed to MAC are aborted only on MAC error or arbitration loss.
REQ-003	6.5.2, 6.5.3	P1	Retransmission 1) Frames losing arbitration shall be retransmitted unless aborted by the LLC user. 2) Non-XL nodes shall support unlimited retransmissions.	
REQ-004	6.5.4	P3	Frame acceptance filtering 1) Receivers should use acceptance filtering to decide frame relevance. 2) Filtering should be a single self-contained operation with no state carry-over from previous frames.	
REQ-005	6.5.6	P1	LLC transfer status reporting LLC shall report five transfer status values: Ongoing, Transmitted, Lost_Arbitration, Disturbed, and Aborted, each reflecting the corresponding MAC sub-layer event.	
REQ-006	6.6.4.4	P1	CRC rules 1) CRC_15 for CC frames. CRC_17 for FD frames with payload ≤16 B. CRC_21 for FD frames with payload >16 B. 2) Init vector: all-zero for CRC_15. MSB-one for CRC_17 and CRC_21.	

ID	ISO ref	Priority	Requirement	Notes
REQ-007	6.6.5.2, 6.6.5.3, 6.6.6.2, 6.6.6.3	P1	Error and overload flags 1) Active error and overload flags: 6 consecutive dominant bits each. 2) Passive error flag: 6 consecutive recessive bits (may be overwritten by dominant bits from other nodes). 3) Error and overload delimiters: 8 recessive bits each. 4) After the flag, all nodes monitor the bus and simultaneously begin the remaining 7-bit delimiter on detecting the first recessive bit.	
REQ-008	6.6.7.1, 6.6.7.2, 6.6.7.3, 6.6.7.4	P1	Inter-frame space 1) Intermission shall be exactly 3 recessive bits. 2) No DF or RF may be started during intermission. 3) EFs and OFs shall not be preceded by inter-frame space. 4) Bus idle is recognised at the third recessive intermission bit (error-active nodes/receivers), the eighth recessive suspend bit (error-passive transmitters), or on leaving bus-integration state. 5) An error-passive transmitter shall suspend for 8 bit times after intermission. If another node transmits during the suspend window, the node shall become a receiver.	
REQ-009	6.6.7.5	P1	Bus re-integration 1) Nodes shall start bus re-integration on protocol startup or bus-off recovery. 2) Re-integration completes after 11 consecutive recessive bits at the sample point. 3) For FD-capable nodes, any synchronisation-causing edge shall reset the 11-bit recessive-bit counter.	
REQ-010	6.6.7.3, 6.6.8	P1	Frame initiation 1) SOF shall be a single dominant bit. 2) A node shall send SOF only when the bus is idle. 3) A dominant third intermission bit causes a node with a pending frame (if error-active or previous receiver) to skip SOF and begin arbitration immediately. 4) A dominant bit detected during bus idle shall be interpreted as SOF.	

ID	ISO ref	Priority	Requirement	Notes
REQ-011	6.6.10.1, 6.6.10.4	P2	Remote frame 1) RTR bit shall be dominant for DFs and recessive for RFs. 2) RFs have no data field. DLC indicates the requested data length.	
REQ-012	6.6.10.2, 6.6.11.2	P1	SRR and RRS reserved bits 1) SRR (CE, FE): transmitted recessive. Receivers accept either polarity without form error. 2) RRS (FB, FE): transmitted dominant. Receivers accept either polarity without form error.	
REQ-013	6.6.10.5, 6.6.11.5	P1	CRC bit stream coverage 1) CC: SOF, arbitration, control, and data fields (dynamic stuff bits excluded). 2) FD: SOF, arbitration, control, data, dynamic stuff bits, and stuff count field (fixed stuff bits excluded).	
REQ-014	6.6.10.6, 6.6.11.6, 6.6.14	P1	Frame acknowledgement 1) Receivers shall ACK consistent frames and flag inconsistent frames with an EF. 2) Transmitters shall flag unacknowledged frames with an EF. 3) FD frames tolerate up to a 2-bit dominant ACK overlap to compensate for receiver phase differences. 4) The ACK delimiter shall be one recessive bit. 5) ACK shall not depend on the frame identifier.	
REQ-015	6.6.10.3, 6.6.11.2, 6.6.11.3	P1	CC/FD format differentiation 1) In FD frames the FDF bit is recessive. The overlapping CC bit position (r0 in CBFF, r1 in CEFF) shall be dominant. 2) IDE is in the control field for CB/FB frames, and in the arbitration field for CE/FE frames. 3) The res bit (FB/FE) shall be dominant.	
REQ-016	6.6.11.3	P2	ESI bit transmission 1) Recessive when the LLC ESI flag is set. 2) Dominant when LLC ESI flag is clear and the node is error-active. 3) Recessive when LLC ESI flag is clear and the node is error-passive.	

ID	ISO ref	Priority	Requirement	Notes
REQ-017	6.6.11.5	P1	Stuff bit count (SBC) field 1) SBC: 3-bit Gray-coded count of dynamic stuff bits (0-7) plus parity bit, placed before the CRC field. 2) TX shall transmit its dynamic stuff-bit count (up to the first FSB) as the SBC field. 3) RX shall verify the received SBC against its own count.	
REQ-018	6.6.10.5, 6.6.11.5	P1	CRC delimiter 1) CC frames: exactly one recessive CRC delimiter bit. 2) FD frames: transmitter sends one recessive bit. A second recessive bit is tolerated before the ACK slot.	
REQ-019	6.6.13.1, 6.6.13.2, 6.6.13.3.1	P1	Bit stuffing 1) After 5 consecutive identical bits, insert a complementary dynamic stuff bit. RX discards it. 2) CC: dynamic stuffing spans SOF to FCRC. FD: SOF to end of data field. ACK and EOF are not stuffed. 3) FD CRC field: a fixed stuff bit precedes the SBC field and follows every fourth CRC bit, each the inverse of the preceding bit. No double stuffing. 4) RX discards FSBs and flags a form error if an FSB matches the preceding bit.	
REQ-020	6.6.10.7, 6.6.11.7, 6.6.15.2	P1	End of frame (EOF) 1) EOF shall consist of exactly 7 recessive bits. 2) A receiver shall validate a frame after the sixth EOF bit. 3) A transmitter requires all seven EOF bits to be error-free.	

ID	ISO ref	Priority	Requirement	Notes
REQ-021	6.6.10.2, 6.6.11.2, 6.6.17.4, 6.6.17.5	P1	Bus arbitration 1) Recessive-sent, dominant-monitored: lose arbitration and become a receiver. 2) Arbitration field: CB: base ID + RTR. CE: base ID + SRR + IDE + extended ID + RTR. FB: base ID + RRS. FE: base ID + SRR + IDE + extended ID. 3) CB wins over FB: CBFF sends dominant r0 where FBFF sends recessive FDF. 4) CB/FB wins over CE/FE: CEFF/FEFF sends recessive SRR where CBFF/FBFF sends dominant. 5) (CB, CE only) DF wins over RF with same ID via its dominant RTR bit.	
REQ-022	6.6.21.2	P1	MAC error detection 1) Bit error: sent/monitored mismatch. Exceptions: recessive non-stuff bit during arbitration, recessive bit in ACK slot, dominant bit while sending a passive error flag. 2) Stuff error: sixth consecutive identical bit in a dynamically-stuffed field. 3) Form error: unexpected fixed-form bit value, or FSB not inverse of preceding bit (FB/FE). Exception: dominant last EOF bit or last delimiter bit is not a form error. 4) CRC error: CRC mismatch. FD also requires SBC match and correct parity. 5) ACK error: transmitter detects no dominant bit in the ACK slot.	
REQ-023	6.6.21.3.1	P1	Error flag initiation and CRC error behaviour 1) On any detected error, start an EF at the next bit and inform the LLC. 2) Data-phase errors: switch to nominal bit rate before starting the EF. 3) CRC error: receiver withholds ACK and sends an EF at the first EOF bit (CC) or 3 bit-times after the CRC delimiter (FD).	

ID	ISO ref	Priority	Requirement	Notes
REQ-024	6.6.21.3.2	P1	Overload frame 1) LLC-requested OF: optional, starts at the first intermission bit, at most 2 per inter-frame space. 2) Reactive OF: triggered by a dominant bit at the first two intermission bits, the last EOF bit (receivers only), or the last error/overload delimiter bit. Starts one bit after the trigger.	
REQ-025	7.3.2, 7.3.3	P1	PCS bit timing configuration 1) Separate nominal and FD data bit rate configurations, each based on a programmable integer prescaler. 2) Sample point at end of Phase_Seg1. Sync_Seg = 1 tq. 3) IPT ≤ 2 tq.	Deployment constraints on the system (not RTL behaviors): prop_seg shall be programmed $\geq t_{\text{node(A)}} + t_{\text{node(B)}} + 2 \times t_{\text{busline}}$ (§7.3.2). SJW shall not exceed min(Phase_Seg1, Phase_Seg2) in either phase (§7.3.2). FD data phase: Phase_Seg2 shall be programmed $\geq$ SJW and $\geq$ maximum IPT (§7.3.3). Oscillator tolerance df shall satisfy $df < SJW / (20 \times t_{\text{bit}})$ and $df < \min(\text{Phase_Seg1}, \text{Phase_Seg2}) / (2 \times (13 \times t_{\text{bit}} - \text{Phase_Seg2}))$ (§7.3.6).
REQ-026	7.3.4, 6.6.21.3.1	P1	Transmitter delay compensation 1) TDC applies in the FD data phase when BRS is recessive. Enable and prescaler (1 or 2) are programmable. 2) SSP position = measured transmitter delay + configured offset. 3) TDC active: bit errors at the SP are ignored. SSP errors are flagged at the following SP.	

ID	ISO ref	Priority	Requirement	Notes
REQ-027	7.3.5.1, 7.3.5.2, 7.3.5.3, 7.3.5.4	P1	PCS synchronisation 1) One synchronisation per bit time, re-enabled after the next recessive SP. 2) Edge qualifies only if the previous SP was recessive and no positive phase error exists while sending dominant. 3) Hard sync: IFS edges (except first intermission bit), bus-integration, and FDF-to-res transition. Restarts bit time with Sync_Seg completed, unconstrained by SJW. 4) All other qualifying edges cause resync, except for the transmitter during the FD data phase. Positive error: Phase_Seg1 $\neq$ SJW. Negative error: Phase_Seg2 $\neq$ SJW. Error $\leq$ SJW: same effect as hard sync.	RTL is stricter than ISO: mac_i.transmitting suppresses all sync unconditionally (line 177), not just resync on positive phase error. Safe since the transmitter is the timing source.
REQ-028	6.6.7.5, 7.4.2.2, 7.4.2.3, 8.1.4.4	P1	PCS bus-off isolation and signalling 1) During bus-off, the PCS shall not drive dominant bits. 2) PCS enters bus-off on a supervisor request and exits on a release request.	
REQ-029	8.1.3.2, Table 14, Table 15	P1	FCE reset and response 1) TEC and REC shall both be reset to zero when the FCE enters its initial state.	
REQ-030	8.1.4.2 rule a), rule b), rule c), rule d), rule e), rule f), rule g), rule h)	P1	Error counter rules 1) REC +1 on any receiver-detected error (except bit errors during error/overload flag). 2) REC +8: first dominant bit after sending an error flag, or bit error while sending an active error/overload flag. 3) REC -1 on successful RX if 1-127. Stays 0 if already 0. Set to 119-127 if above 127. 4) TEC +8 when sending an error flag. Exceptions: error-passive ACK error, pre-arbitration stuff error. 5) TEC +8 on bit error while sending an active error or overload flag. 6) TEC -1 on successful TX (floor 0). 7) After 14 consecutive dominant bits post active error/overload flag (8 post passive error flag): both +8. Each additional 8 dominant bits: both +8.	

ID	ISO ref	Priority	Requirement	Notes
REQ-031	8.1.4.3, 8.1.4.4, 8.1.3.2	P1	Error confinement state transitions 1) Error-passive when either error counter > 127. Error-active restored when both ≤ 127. 2) TEC > 255 causes bus-off. 3) Bus-off recovery requires 128 idle conditions. Both counters reset to zero on recovery. 4) FCE shall notify the LLC on entering bus-off state.	Sub-claim 4: bus-off recovery requires 128 separate idle detections (128 × 11 recessive bits), not 128 consecutive recessive bits.
REQ-032	6.6.11.3	P1	BRS bit rate switching 1) Recessive BRS: switch to FD data rate at the BRS sample point. Dominant BRS: keep nominal rate. 2) Rate switches back to nominal at the first CRC delimiter sample point.	
REQ-033	6.4.3, 6.4.4, Table 5	P1	DLC byte-count mapping DLC 0-8: linear mapping. FD: DLC 9→12, 10→16, 11→20, 12→24, 13→32, 14→48, 15→64 bytes. CC: DLC 9-15 clamped to 8 bytes.	
REQ-034	6.6.16	P1	Byte/Bit transmission order Within the data field byte 0 is sent first. Within each byte bit(7) is sent first.	
REQ-035	8.1.5	P1	Node start-up 1) A node shall complete bus re-integration (11 consecutive recessive bits) before transmitting. 2) A single node receiving no ACK becomes error-passive but never bus-off.	
REQ-036	6.6.18	P2	MAC data consistency Frames shall be fully buffered before TX starts, or the LLC shall detect and invalidate frames corrupted mid-transmission.	Verify by code inspection which method the implementation uses and confirm it is applied consistently. If method (a) - confirm the MAC FSM only begins serialisation after the full frame is loaded. If method (b) - confirm the LLC has an error-detection path and can abort or invalidate an in-progress frame.

ID	ISO ref	Priority	Requirement	Notes
REQ-037	6.6.21.1, 6.6.21.4	P2	Error signalling enable A configuration parameter shall exist to enable or disable error signalling.	
REQ-038	6.4.4	P3	DLC padding An LLC restricted to fewer bytes shall replace excess bytes with 0xCC padding in both TX and RX directions.	Only applicable if the implementation exposes a configurable maximum-data-byte restriction. If not implemented, waive.

A.2 Verification Plan Fields

The verification plan table contains the verification metadata fields added during the planning phase. Cross-reference to the requirements table above via the shared ID field.

Table 9: ISO 11898-1 verification plan.

ID	Layer	Side	Format	Observability	Method	Status	Label	File	Coverage criteria
REQ-001	LLC	both	CB, CE, FB, FE	black_box	simulation	not_started			
REQ-002	LLC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started			
REQ-003	LLC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started			
REQ-004	LLC	receiver	CB, CE, FB, FE	white_box	code_inspection	not_started			
REQ-005	LLC	transmitter	CB, CE, FB, FE	black_box	simulation	not_started			
REQ-006	MAC	both	CB, CE, FB, FE	white_box	simulation, coverage	in_progress	crc_15, crc_17, crc_21	can_mac_crc_tb	Three bins: (1) CB or CE frame (CRC_15), (2) FB or FE frame with data ≤ 16 B (DLC ≤ 10 , CRC_17), (3) FB or FE frame with data > 16 B (DLC ≥ 11 , CRC_21) - must each be hit.
REQ-007	system	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm:~900, test_bus_off	can_mac_pcs_fce_tb	
REQ-008	MAC	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm:~486, test_normal, test_bus_off	can_mac_pcs_fce_tb	
REQ-009	system	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm:~433, test_bus_off	can_mac_pcs_fce_tb	

ID	Layer	Side	Format	Observability	Method	Status	Label	File	Coverage criteria
REQ-010	MAC	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~451, test_normal	can_mac_pcs_fce_tb	
REQ-011	MAC	both	CB, CE	black_box	simulation	not_started			
REQ-012	MAC	both	CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~546, p_fsm::~548, test_normal	can_mac_pcs_fce_tb	
REQ-013	MAC	both	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm::~299, p_fsm::~216	can_mac_fsm	
REQ-014	system	both	CB, CE, FB, FE	white_box	simulation, wave-form_inspection	in_progress	test_normal	can_mac_pcs_fce_tb	
REQ-015	MAC	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~613, test_normal	can_mac_pcs_fce_tb	
REQ-016	MAC	transmitter	FB, FE	white_box	simulation, wave-form_inspection	in_progress	test_normal	can_mac_pcs_fce_tb	
REQ-017	MAC	both	FB, FE	black_box	simulation	in_progress	p_sbc_checker	can_mac_bs_tb	
REQ-018	MAC	transmitter	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~768, p_fsm::~796, test_normal	can_mac_pcs_fce_tb	
REQ-019	MAC	both	CB, CE, FB, FE	white_box	simulation, coverage	in_progress	Input Coverage, Output Coverage	can_mac_bs_tb	All bins in cov_input and cov_output must be hit.
REQ-020	MAC	both	CB, CE, FB, FE	white_box	simulation	in_progress	test_normal	can_mac_pcs_fce_tb	
REQ-021	system	transmitter	CB, CE, FB, FE	black_box	simulation	in_progress	test_lost_arb	can_mac_pcs_fce_tb	

ID	Layer	Side	Format	Observability	Method	Status	Label	File	Coverage criteria
REQ-022	MAC	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~283, p_fsm::~336, p_fsm::~343, p_fsm::~415, test_bus_off	can_mac_pcs_fce_tb	
REQ-023	MAC	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_fsm::~283, p_fsm::~294, p_fsm::~357, test_bus_off	can_mac_pcs_fce_tb	
REQ-024	MAC	both	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm::~370	can_mac_fsm	
REQ-025	PCS	both	CB, CE, FB, FE	white_box	simulation	in_progress	test_normal	can_pcs_tb	
REQ-026	PCS	transmitter	FB, FE	black_box	simulation	in_progress	p_check_tdc_delay	can_pcs_tb	
REQ-027	PCS	both	CB, CE, FB, FE	white_box	code_inspection, wave-form_inspection	in_progress	p_can_pcs::~177, p_can_pcs::~228, test_normal	can_pcs_tb	
REQ-028	PCS	both	CB, CE, FB, FE	white_box	simulation, code_inspection	in_progress	p_can_pcs::~280, p_can_pcs::~379, test_bus_off	can_mac_pcs_fce_tb	
REQ-029	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_reset, test_bus_off_recovery	can_fce_tb	Sub-claim 1: TEC/REC reset shall take effect within one clock cycle of the request. Sub-claim 2: verify that Normal_mode_response is asserted after reset completes.

ID	Layer	Side	Format	Observability	Method	Status	Label	File	Coverage criteria
REQ-030	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_rule_a, test_rule_b, test_rule_c, test_rule_c_exception, test_rule_d, test_rule_e, test_rule_f_tx, test_rule_f_rx, test_rule_g, test_rule_h	can_fce_tb	Rule h: the '>127 → 119-127' clamp path must be explicitly confirmed in simulation.
REQ-031	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_bus_off, test_bus_off_recovery	can_fce_tb	
REQ-032	MAC	both	FB, FE	white_box	code_inspection, wave- form_inspection	in_progress	p_fsm::~636, p_fsm::~755, test_normal	can_mac_pcs_fce_tb	
REQ-033	LLC	both	CB, CE, FB, FE	black_box	simulation, coverage	not_started			Four non-linear FD DLC values (9, 12, 14, 15 → 12, 24, 48, 64 bytes) plus DLC=8 boundary must each produce a correctly sized frame.
REQ-034	MAC	both	CB, CE, FB, FE	black_box	simulation	not_started			

ID	Layer	Side	Format	Observability	Method	Status	Label	File	Coverage criteria
REQ-035	system	both	CB, CE, FB, FE	black_box	simulation	not_started			Sub-claim 1: confirm no TX before 11 consecutive recessive bits are detected at the SP. Sub-claim 2: single-node scenario with no ACK injection; TEC must reach error-passive but not exceed 255.
REQ-036	MAC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started			
REQ-037	MAC	both	CB, CE, FB, FE	white_box	code_inspection	not_started			
REQ-038	LLC	both	CB, CE, FB, FE	black_box	simulation	not_started			

- [1] International Organization for Standardization, *Road vehicles — controller area network (CAN) — Part 1: Data link layer and physical coding sublayer*. 2024.
- [2] M. Jeřábek and P. Píša, “CTU CAN FD — open-source CAN FD IP core.” 2019. Available: https://gitlab.fel.cvut.cz/canbus/ctucanfd_ip_core
- [3] M. Jeřábek, “Open-source and open-hardware CAN FD protocol support,” PhD thesis, Czech Technical University in Prague, 2019. Available: <https://dspace.cvut.cz/handle/10467/82630>
- [4] International Organization for Standardization, *Road vehicles — controller area network (CAN) conformance test plan — Part 1: Data link layer and physical signalling*. 2016.
- [5] OpenCores Contributors, “CAN protocol controller — OpenCores.” 2003. Available: <https://opencores.org/projects/can>
- [6] S. V. Nøsen, “Canola — CAN controller in VHDL.” 2020. Available: <https://github.com/svnesbo/canola>
- [7] Robert Bosch GmbH, “M_CAN — CAN FD protocol controller IP module.” 2024. Available: <https://www.bosch-semiconductors.com/ip-modules/can-fd/>
- [8] F. Hartwich, “CAN with flexible data-rate,” in *Proceedings of the 13th international CAN conference (iCC 2012)*, 2012.
- [9] AMD (Xilinx), “CAN FD controller — Vivado design suite product guide (PG223).” 2024. Available: <https://docs.amd.com/r/en-US/pg223-canfd>
- [10] CAST, Inc., “CAN-FD protocol controller.” 2024. Available: <https://www.cast-inc.com/automotive/can/can-fd-protocol-controller>
- [11] GHDL Contributors, “GHDL — open-source VHDL simulator.” 2024. Available: <https://ghdl.github.io/ghdl/>
- [12] OSVVM Contributors, “OSVVM — open source VHDL verification methodology.” 2024. Available: <https://osvvm.org>
- [13] Intel Corporation, “Avalon interface specifications,” Intel Corporation, 683091, 2021. Available: <https://www.intel.com/programmable/technical-pdfs/683091.pdf>
- [14] J. Bergeron, “Verification planning,” in *Writing testbenches: Functional verification of HDL models*, 2nd ed., Kluwer Academic Publishers, 2003.