

Implementation and Verification of a CAN-FD Bus Transceiver in VHDL

Mads Richardt (s224948)

February 28, 2026

Contents

Abstract	2
Abbreviations	3
1 Introduction	3
1.1 Motivation	4
1.2 Existing CAN Controller	4
1.2.1 Limitations	4
1.2.2 Decision to Redesign	5
1.3 Existing CAN FD IP Cores	5
1.3.1 Open-Source Implementations	5
1.3.2 Commercial Implementations	5
1.3.3 Rationale for In-House Development	6
1.4 Problem Statement	7
1.5 Objectives	7
2 Background	7
2.1 CAN Protocol Evolution	7
2.2 ISO 11898-1:2024 Standard	7
2.3 VHDL and OSVVM	8
3 Requirements Engineering & Verification Planning	8
3.1 VHDL Code Standard and Design Constraints	8
3.1.1 Project-specific Infrastructure Requirements	8
3.1.2 File Structure and Naming	9
3.1.3 RTL Coding Style and Verification	9
3.2 From Standard to Structured Requirements	9
3.2.1 Layer	10
3.2.2 Side	10
3.2.3 Format Applicability	10
3.2.4 Observability	11
3.2.5 Priority	11
3.3 Verification Plan Data Structure	11
3.3.1 Verification Method	11
3.3.2 Traceability: Label and File	11
3.3.3 Status	11
3.4 Storage Format and Tooling	12
3.5 AI-Assisted Extraction: Utility and Limitations	13

4	Design and Architecture	13
4.1	Ramifications of the Requirements Model on Initial Design Strategy	13
4.2	Architectural Design Decisions	14
4.2.1	Monolithic vs. Layered Architecture	14
4.2.2	Combined vs. Separated TX/RX FSMs	15
4.2.3	Per-Field vs. Per-Phase FSM Granularity	16
4.2.4	Interface Record Design	16
4.2.5	Dual CRC Data Feeds	16
4.3	System Overview	17
4.4	LLC Frame Format	17
4.5	LLC Sub-layer	17
4.5.1	can_llc_tx	20
4.5.2	can_llc_rx	20
4.6	MAC Sub-layer	20
4.6.1	can_mac_tx	20
4.6.1.1	can_mac_fsm_tx	20
4.6.1.2	can_mac_ser_tx	22
4.6.1.3	can_mac_bs	22
4.6.1.4	can_mac_crc	25
4.6.2	can_mac_rx	25
4.6.3	can_mac Unified Wrapper	27
4.7	FCE Sub-layer	27
4.8	PCS Sub-layer	27
4.8.1	can_pcs_tx	27
4.8.2	can_pcs_rx	27
5	Implementation	31
5.1	Type Safety and Packages	31
5.2	Bit Timing and TDC	31
5.3	CRC and Bit Stuffing	31
6	Verification and Results	31
6.1	Testbench Results Summary	31
7	Discussion	32
7.1	Future Work	32
8	Conclusion	32
9	References	32
A	Verification Plan	32
A.1	Extracted Requirements	32
A.2	Verification Plan Fields	39

Abstract

*This document is a work-in-progress draft for the B.Eng thesis. Currently, the project is in the **Verification Plan** phase (Phase 2), with architectural prototyping for the transmitter sub-system completed.*

This thesis describes the design, implementation, and verification of a CAN (Controller Area Network) and CAN-FD (Flexible Data rate) transmitter sub-system. The implementation is targeting high-reliability engine controller applications and is compliant with the ISO 11898-1:2024 standard [1]. Key features include

support for both Classic and FD frame formats, Transmitter Delay Compensation (TDC) for high-speed data phases, and a modular architecture separated into Link Layer Control (LLC), Media Access Control (MAC), and Physical Coding Sublayer (PCS).

Abbreviations

Table 1: Abbreviations used in this report.

Abbreviation	Meaning
ACK	Acknowledgement
AF	Acceptance Field
AUI	Attachment Unit Interface
CAN	Controller Area Network
CBFF	Classical Base Frame Format
CEFF	Classical Extended Frame Format
DF	Data Frame
DLL	Data Link Layer
EF	Error Frame
FBFF	FD Base Frame Format
FCE	Fault Confinement Entity
FEFF	FD Extended Frame Format
FTYP	Frame Type
LLC	Logical Link Control
LSDU	LLC Service Data Unit
MAC	Medium Access Control
OF	Overload Frame
OSI	Open Systems Interconnection
PCS	Physical Coding Sublayer
PDU	Protocol Data Unit
PMA	Physical Medium Attachment
RF	Remote Frame
SAP	Service Access Point
SDU	Service Data Unit
SOF	Start of Frame
TEC/REC	Transmit Error Counter / Receive Error Counter
VCID	Virtual CAN Channel Identifier

1 Introduction

TODO: Add a section on the IO extender board (The board is on the test wall, get the name from Alex)

1.1 Motivation

The Controller Area Network (CAN) has been the workhorse of automotive and industrial communication for decades. However, the increasing bandwidth requirements of modern systems led to the development of CAN-FD (Flexible Data rate), which allows for larger payloads and higher bit rates. This project aims to provide a robust, hardware-independent VHDL implementation of a CAN-FD transmitter.

1.2 Existing CAN Controller

The starting point for this project is an existing CAN Classic controller developed internally at the company. The controller is implemented in VHDL and has been integrated into a production IO-extender FPGA design. It supports CAN Classic frames with both 11-bit (base) and 29-bit (extended) identifiers at bit rates up to 500 kbit/s, and has been verified through hardware bring-up on physical CAN buses.

The controller follows a monolithic architecture where a single top-level wrapper (`can_bus_controller`) instantiates six sub-modules: a combined TX/RX frame FSM (`can_fsm`), a bit timing generator (`can_node_clock`), a dynamic bit stuffer (`can_stuff_bit_gen`), a CRC-15 engine (`gen_crc`), and two Avalon-ST converters for frame serialization and deserialization (`can_ast_to_serial`, `can_serial_to_ast`). The entire design totals roughly 2,000 lines of VHDL across seven source files.

The central component is `can_fsm`, an 810-line state machine with 18 states that handles both transmission and reception in a single process. The FSM manages frame arbitration, bit-level transmission and reception, stuff bit error checking, CRC validation, ACK handling, and error flag generation. Error counting (TEC/REC) and node state transitions (Error Active, Error Passive, Bus Off) are handled in separate processes within the same file but are tightly coupled to the main FSM through shared signals.

1.2.1 Limitations

While the existing controller is functional for CAN Classic, several architectural limitations prevent it from being extended to support CAN FD:

Single bit rate domain. The `can_node_clock` module generates a single pair of timing strobes - one sample pulse and one transmit pulse - derived from a fixed set of bit timing parameters. CAN FD requires switching between a nominal bit rate (used during arbitration) and a faster data bit rate (used during the data phase), with Transmitter Delay Compensation (TDC) to account for the transceiver round-trip delay at the higher rate. Adding dual bit rate support and TDC to the existing clock module would require a fundamental redesign of the timing architecture.

Dynamic bit stuffing only. The `can_stuff_bit_gen` module implements the CAN Classic rule of inserting an inverse bit after five consecutive identical bits. CAN FD introduces a second stuffing mode - fixed bit stuffing - where a stuff bit is inserted at fixed intervals during the CRC field, and a Stuff Bit Count (SBC) field with Gray-coded parity is appended. The existing stuffer has no mechanism for mode switching or SBC generation.

Single CRC polynomial. The controller uses a single CRC-15 instance. CAN FD requires three CRC polynomials: CRC-15 for Classic frames, CRC-17 for FD frames with payloads up to 16 bytes, and CRC-21 for larger FD payloads. Furthermore, the CRC data feed differs between Classic and FD - in FD frames, dynamic stuff bits in the arbitration region are included in the CRC computation, requiring a dual data feed to the CRC engine.

Combined TX/RX FSM. The monolithic FSM interleaves transmission and reception logic in every state,

with the `is_transmitter` flag selecting the active code path. This coupling makes it difficult to add FD-specific states (such as BRS, ESI, and the SBC field) without increasing the already high cyclomatic complexity. A CAN FD frame has more control fields than a Classic frame, and handling both TX and RX paths for all four frame formats (Classic Base, Classic Extended, FD Base, FD Extended) in a single process would result in an unwieldy state machine.

Embedded error handling. Error detection, error flag transmission, and error counter management are distributed across the main FSM process and its auxiliary processes. The ISO 11898-1 standard defines the Fault Confinement Entity (FCE) as a logically separate component with well-defined interfaces to the MAC and PCS sub-layers. Extracting the error handling into a reusable, independently testable FCE module - as required by the standard's layered architecture - would require significant refactoring of the existing FSM.

1.2.2 Decision to Redesign

Given these limitations, extending the existing controller to CAN FD would require modifying nearly every sub-module and fundamentally restructuring the FSM. The resulting design would carry the constraints of the original monolithic architecture while trying to support a significantly more complex protocol. Instead, a clean-sheet redesign was chosen, structured around the ISO 11898-1 layered architecture (LLC, MAC, PCS, FCE) with separated TX and RX paths, reusable sub-components (bit stuffer, CRC engine), and typed record interfaces. This approach allows each sub-layer to be implemented and verified independently, supports both Classic and FD frame formats from the outset, and produces a modular design that can be maintained and extended without cross-cutting changes.

1.3 Existing CAN FD IP Cores

Before committing to an in-house redesign, the available CAN FD controller IP cores were evaluated. Table 2 summarizes the candidates, spanning both open-source and commercial offerings.

1.3.1 Open-Source Implementations

CTU CAN FD [2], [3] is the only mature open-source CAN FD controller available as synthesizable HDL. Developed at the Czech Technical University in Prague, it is written in VHDL, licensed under MIT, and has been conformance-tested against ISO 16845-1 [4]. The controller includes a full TX and RX pipeline with up to four TX buffers, acceptance filtering, timestamping, and a register interface with DMA support. A mainline Linux kernel driver has been available since kernel version 5.12. CTU CAN FD represents a complete, production-oriented CAN node - a significantly broader scope than what is needed in this project.

OpenCores CAN [5] is a Verilog controller modeled after the Philips SJA1000 register interface. It is one of the earliest open-source CAN cores and is widely cited in academic work. However, it supports only CAN 2.0B (Classic CAN) and has seen no active development since approximately 2010. It cannot serve as a starting point for CAN FD.

Canola [6] is a VHDL CAN 2.0B controller with a clean VHDL-2008 codebase and a cocotb-based testbench. It includes a Triple Modular Redundancy (TMR) wrapper for radiation-tolerant applications. Like the OpenCores core, it does not support CAN FD.

1.3.2 Commercial Implementations

Bosch M_CAN [7], [8] is the reference CAN FD controller, developed by the inventor of both CAN and CAN FD. M_CAN is the IP core embedded in virtually every automotive microcontroller (NXP S32, Infineon

AURIX, STM32, TI Jacinto, Renesas RH850). It is licensed under NDA with per-design royalty fees.

AMD/Xilinx CAN FD [9] is a soft IP core included in the Vivado Design Suite. It provides an AXI4-Lite register interface with up to 32 acceptance filters, TX mailboxes, and RX FIFOs. It is device-locked to AMD/Xilinx FPGAs and cannot be ported to other targets.

CAST CAN FD [10] is a technology-independent RTL core with AMBA APB/AHB bus interface options. It is licensed per-design with an upfront fee. Synopsys (DesignWare) and Cadence offer similar ASIC-targeted CAN FD cores under their respective IP licensing programs.

Table 2: Survey of available CAN FD controller IP cores.

Implementation	Language	CAN FD	License	Scope	Conformance Tested
CTU CAN FD [2]	VHDL	Yes	MIT	Full node (TX+RX, buffers, DMA)	ISO 16845-1
OpenCores CAN [5]	Verilog	No	LGPL	Full node (CAN 2.0B only)	No
Canola [6]	VHDL	No	MIT	Full node (CAN 2.0B, TMR)	No
Bosch M_CAN [7]	HDL (NDA)	Yes	Per-design royalty	Full node	Yes (reference)
AMD/Xilinx CAN FD [9]	HDL	Yes	Vivado-included	Full node	Yes
CAST CAN FD [10]	RTL	Yes	Per-design fee	Full node	Yes

1.3.3 Rationale for In-House Development

Despite the availability of these solutions, none satisfies the combined requirements of a large engine design company developing safety-critical control systems. The decision to develop the CAN FD transceiver in-house was driven by five considerations.

TODO: Add something about the 30 year maintenance requirements that the company is responsible for (hard to rely on 3rd party products for this long). **IP ownership and supply chain independence.** Commercial IP cores introduce a licensing dependency on an external vendor. For a product with a multi-decade lifecycle - typical in heavy-duty engine applications - this dependency carries risk: vendors may discontinue support, change licensing terms, or be acquired. Owning the RTL outright eliminates this exposure and ensures that the design can be maintained, ported, and modified without third-party approval for the full product lifetime. The open-source CTU CAN FD avoids the licensing risk, but using it still means adopting a codebase whose architecture, naming conventions, and design decisions were made for a different context.

Verification authority. In safety-critical domains, the verification evidence must be traceable from standard requirements to RTL assertions and testbench results. Adopting a third-party core - even one conformance-tested against ISO 16845-1 - means inheriting its verification artifacts rather than producing them. The company's verification methodology requires full control over the verification plan, the testbench architecture, and the assertion coverage. Building the RTL in-house allows the verification plan

(described in Section ??) to drive the implementation, ensuring that every module is verified against the specific requirements extracted from the standard, using the company's own toolchain and conventions.

Architectural scope. All available IP cores implement a complete CAN node: TX and RX pipelines, message RAM, acceptance filtering, buffer management, register interfaces, and in some cases DMA controllers. The company's application requires only the data link layer (LLC, MAC, FCE) and physical coding sublayer (PCS) - the protocol engine that converts between byte-level frame data and the serial bus. The higher-level buffering and filtering logic already exists in the company's FPGA infrastructure. Adopting a full-node IP core would introduce unnecessary complexity and area overhead, and the integration effort to bypass or disable the unused subsystems may approach the effort of a targeted implementation.

Integration with existing infrastructure. The company's FPGA designs use a specific Avalon-ST streaming interface for inter-module communication, a particular clock and reset architecture, and established conventions for signal naming and module boundaries. A third-party core would require an adaptation layer to bridge its native interface (AXI, APB, or custom register map) to the existing infrastructure. The in-house design uses the company's interface conventions natively, eliminating this integration overhead.

Platform independence. The AMD/Xilinx CAN FD core is locked to Xilinx devices. The Bosch M_CAN and other commercial cores are delivered as technology-specific netlists or encrypted RTL for a particular target. The in-house design is written in portable VHDL-2008, synthesizable on any FPGA platform or ASIC flow, ensuring that the IP remains usable if the company changes FPGA vendors.

1.4 Problem Statement

The need for CAN FD support in the company's engine controller platform, combined with the architectural limitations of the existing CAN Classic controller (Section 1.2.1) and the unsuitability of available third-party IP cores for the company's specific requirements (Section 1.3.3), motivates the development of a new CAN FD transceiver from the ground up. The challenge lies in creating a modular, independently verifiable implementation that supports both Classic and FD frame formats, handles dual bit rate switching with Transmitter Delay Compensation (TDC), and maintains strict compliance with ISO 11898-1 [1] - all while fitting into the company's existing FPGA infrastructure and verification methodology.

1.5 Objectives

- Implement a VHDL-2008 compliant CAN/CAN-FD transmitter.
- Support both Base (11-bit) and Extended (29-bit) identifiers.
- Implement TDC measurement and compensation logic.
- Ensure high verification coverage using OSVVM and GHDL.

2 Background

2.1 CAN Protocol Evolution

Brief history from CAN 2.0 to CAN-FD.

2.2 ISO 11898-1:2024 Standard

Overview of the data link layer and physical signaling requirements [1].




Figure 1: CAN-Bus consisting of three CAN nodes.

2.3 VHDL and OSVVM

The role of modern VHDL standards and verification frameworks in digital design.

3 Requirements Engineering & Verification Planning

Bridging the gap between a normative specification document and a traceable verification plan is a pivotal task in any protocol implementation project. The present section describes the process of distilling a coherent set of verifiable requirements and associated verification plan from the ISO 11898-1 standard. In addition, the constraints imposed by general company practices and standards are described along with project-specific requirements related to the larger system in which the module is intended to integrate.

3.1 VHDL Code Standard and Design Constraints

The constraints on this project come from two distinct sources. Two requirements are specific to this project's place within the company's existing CAN infrastructure while the remaining constraints come from the company's VHDL Code Standard and apply uniformly to all FPGA IP modules developed in-house.

3.1.1 Project-specific Infrastructure Requirements

1. **Avalon-ST user interface:** The CAN controller's external interface to the host system must use the Avalon-ST streaming protocol [11] (data, valid, ready, sop, eop). The requirement applies specifically to the boundary between the CAN controller and its user.
2. **IP library CRC block:** The company maintains a reusable, parameterised CRC generator (`gen_crc`) in its IP library. This module must be used for all CRC computations.

3.1.2 File Structure and Naming

1. **Per-module file structure:** Each module must be organized into a fixed directory layout: `hdl_src/` for RTL source files, `hdl_tb/` for test bench files, and `test_case/` for waveform configuration. One entity per VHDL file, with the filename matching the entity name. Every entity requires a dedicated test bench, unless it is instantiated exclusively as a sub-module of a fully tested parent.
2. **Entity port types:** Entity ports are restricted to `std_logic`, `std_logic_vector`, and records or arrays of these types. Modes are restricted to `in` and `out`.
3. **Naming conventions:** A mandatory prefix/suffix scheme applies to all VHDL identifiers: types (`t_`), constants (`c_`), generics (`gc_`), processes (`p_`), functions (`f_`), packages (`pk_`), state variables (`s_`), entity inputs (`_i`), and entity outputs (`_o`).

3.1.3 RTL Coding Style and Verification

1. **RTL design rules:** Synchronous processes must be sensitive to the clock only. Reset must be synchronous and initialize all control registers. FSMs are preferably implemented as single-process designs where all signal assignments are derived from the current state, with an explicit other-state that returns to a known safe state.
2. **Test bench requirements:** Test benches must follow a black-box testing model and test cases must be ordered as: reset tests first, then a normal-usage test, then all remaining tests.

3.2 From Standard to Structured Requirements

The requirements engineering process was aimed at tackling two key objectives:

1. Extracting a clear and actionable set of requirements that could serve as a starting point for the design phase.
2. Establishing a clear, traceable link between the ISO 11898-1 specification and the verification environment.

Both objectives are complicated by the nature of the source material. Normative requirements are distributed across subsections, often restated from different perspectives, and interspersed with explanatory and descriptive text. Accordingly, extracting a precise unambiguous set of requirements from such prose is non-trivial and inherently prone to oversights and misinterpretations. Nonetheless, a structured and precise requirements set is absolutely essential. It serves as both the starting point for the system design phase and as the target of verification effort. [12]

The requirements set was constructed using the AI-assisted pipeline shown in Figure 2. The first step was converting the ISO 11898-1 pdf to Markdown - a format which can be efficiently searched and ingested by LLM models. The resulting Markdown file was then fed to a Claude Sonnet 4.6 LLM agent, which was prompted to extract all normative statements - sentences containing words like “shall”, “should”, “must”, and their corresponding negations.

This process yielded a raw set of 168 normative statements linked to the ISO standard sections from which they were extracted. The normative set was then manually reviewed, consolidated, and distilled into a final set of 45 requirements. In addition to the ISO standard section links provided by the initial LLM agent extraction, each requirement was classified along the following set of dimensions:

- layer, Section 3.2.1
- side, Section 3.2.2




Figure 2: Pipeline generating the `verification_plan.toml` artifact. 1) LLM agent extraction of normative statements from the ISO 11898-1 standard. 2) Manual grouping of related normative statements and requirement distillation. 3) Argumentation with additional requirement labels and verification specific fields.

- format, Section 3.2.3
- priority, Section 3.2.5
- observability, Section 3.2.4

These dimensions, described in the following subsections, should serve to focus the design phase by carving out the natural sub-module boundaries and guide the verification effort by highlighting the relevant layer of abstraction associated with a given requirement.

3.2.1 Layer

The ISO 11898-1 structures the CAN protocol into three functional sub-layers and a layer-spanning Fault Confinement Entity (FCE):

- Logic Link Control (LLC).
- Medium Access Control (MAC).
- Physical Coding Sub-layer (PCS).

Each of these have a dedicated section in the ISO 11898-1 standard and the layer field assigns each requirement to the protocol sub-layer that owns it. This classification determines the verification boundary at which each requirement must be exercised. A fifth label - **system** - was introduced alongside the four protocol layers to classify requirements that are inherently multi-layer or multi-node in character. Some CAN behaviors cannot be attributed to a single layer of a single node. They emerge from interactions between multiple nodes on the bus, or span the layer boundary within a single node. The system label flags these requirements as ones that require either an integrated multi-module test bench or a multi-node simulation environment.

3.2.2 Side

The side field records whether a requirement pertains to the transmitter path, the receiver path, or both roles simultaneously. This dimension reflects the ISO standard's own framing, which frequently specifies transmitter and receiver obligations separately.

3.2.3 Format Applicability

The `format_applicability` field records which of the four in-scope frame formats (CB, CE, FB, FE) each requirement applies to. This field determines which testbench stimulus configurations are required to exercise a given requirement.

3.2.4 Observability

The observability field resolves each requirement as either black-box or white-box, relative to the module boundary of the owning layer:

- **Black-box:** Can be verified purely through the module’s observable port signals.
- **White-box:** Verification requires direct observation of the module’s internal state.

3.2.5 Priority

The priority field classifies each requirement into one of three levels:

- P1 requirements are need-to-have - they must be verified before the design can be considered complete.
- P2 requirements are nice-to-have - they are verified in the normal verification cycle but do not block closure.
- P3 requirements are optional - addressed only if schedule permits.

3.3 Verification Plan Data Structure

The consolidated requirements set described in Section 3.2 provided the foundation for the verification plan development. The verification plan data structure augments each requirement with the dimensions needed to answer not just *what* must be true, but *how* it will be verified, *where* the evidence lives, and *when* verification is complete. This additional structure follows the verification planning methodology described by Bergeron [12], which organizes a verification plan around explicit links between requirements, verification methods, and traceability artifacts. The complete 45-requirement plan is reproduced in Section A. Table 3 lists the fields attached to each requirement entry; the following sections detail the rationale and allowed values for each field.

3.3.1 Verification Method

The `verification_method` field makes the path from requirement to verification artifact explicit and actionable before any testbench is written, giving each requirement a clear route to closure before implementation begins. Four methods are used: simulation (automated assertion or check procedure in a testbench), code inspection (RTL source review), waveform inspection (manual review of simulation output to confirm bit-level timing or signal sequencing), and coverage (OSVVM coverage bins sweeping a value range). Combinations are valid when multiple sub-claims within one requirement each call for a different method.

3.3.2 Traceability: Label and File

Each requirement entry carries two dedicated traceability fields: a `file` field identifying the testbench or RTL source file responsible for covering the requirement, and a `label` field identifying a specific named procedure, assertion, or coverage ID within that file. Together they establish a direct, navigable link from each requirement to its verification artifact, following the requirements-driven traceability methodology described in [12].

3.3.3 Status

The `status` field (`not_started`, `in_progress`, `complete`) records closure state explicitly, allowing partial progress to be tracked without relying on whether the traceability fields happen to be populated. Crucially, gaps in coverage remain immediately visible: requirements with unpopulated traceability fields and

`status = not_started` are structurally obvious in the data file without requiring a separate coverage matrix.

Table 3: Verification-plan metadata fields.

Field	Purpose
<code>id</code>	Sequential identifier REQ-NNN.
<code>source_clause</code>	ISO 11898-1:2015 section reference.
<code>original_wording</code>	Verbatim normative text from the standard.
<code>paraphrase</code>	Concise restatement for report tables and review.
<code>layer</code>	Sub-layer owner: LLC, MAC, PCS, FCE, or system (see Section 3.2.1).
<code>side</code>	Obligation direction: transmitter, receiver, or both (see Section 3.2.2).
<code>format_applicability</code>	Applicable frame formats: CB, CE, FB, FE (see Section 3.2.3).
<code>observability</code>	<code>black_box</code> or <code>white_box</code> (see Section 3.2.4).
<code>verification_method</code>	Method(s) used to verify the requirement (see Section 3.3.1).
<code>priority</code>	P1 (need-to-have), P2 (nice-to-have), or P3 (optional) (see Section 3.2.5).
<code>status</code>	<code>not_started</code> , <code>in_progress</code> , <code>complete</code> , or <code>waived</code> (see Section 3.3.3).
<code>notes</code>	Engineering notes and caveats.
<code>label</code>	Assertion label, TB procedure name, coverage ID, or RTL tag. Comma-separated when multiple procedures cover distinct sub-claims (see Section 3.3.2).
<code>file</code>	Target file - TB for simulation/coverage, RTL for code inspection. Comma-separated when sub-claims span multiple files (see Section 3.3.2).

3.4 Storage Format and Tooling

The verification plan is stored as a TOML file with each entry as a self-contained `[[requirement]]` block with one key per line. TOML was chosen because it is human-readable, maps cleanly onto structured records without a schema file, and can be parsed and manipulated with minimal code using Python’s built-in `tomllib` library - making it straightforward to write validation and query scripts alongside the MCP server.

A **Model Context Protocol (MCP) server** (`mcp_tools/verification_plan_manager.py`) was written to provide a constrained interface for LLM-assisted operations on the plan thus minimizing the risk of silent data corruption caused by un-constrained LLM-agent interactions with the verification plan source.

- **`get_requirement` / `query_requirements`:** retrieve individual entries or filtered subsets by layer, side, status, observability, or whether traceability fields are populated.
- **`update_requirement`:** atomically update a single named field of a single entry, with schema validation before the write is committed.
- **`insert_requirement`:** add a new entry with automatic field initialization and sequential-ID assignment.

- **delete_requirement**: remove an entry and flag any cross-references to it.
- **renumber_requirements**: regenerate all REQ-NNN identifiers after insertions or deletions, maintaining a consistent sequential namespace.
- **get_statistics**: report coverage of label, file, and status fields to give an instant snapshot of plan completeness.

The verification plan TOML operated as a living document throughout the project. As implementation and verification work progressed, traceability links were added as testbenches were written, and status fields were updated to reflect closure. The MCP tooling made these ongoing updates tractable: individual field changes were atomic, sequentially logged, and immediately visible in version-control diffs.

3.5 AI-Assisted Extraction: Utility and Limitations

The LLM agent earned its keep by bootstrapping the initial content of the verification plan data structure. Starting from a blank requirements table would have required a significantly larger upfront effort, and having a populated, classified starting point - even one requiring substantial revision - gave the manual review process a concrete artifact to work from and react to. Generating a rough first draft from source material is a well-understood and widely used application of language models, and this case was no exception.

That said, the overall time saving was marginal. The most time-consuming part of the process - close reading, semantic consolidation, and orthogonality-driven reduction from 168 to 45 entries - is equally demanding whether the starting point is a blank table or an AI-generated draft. The agent's output had to be reviewed statement by statement, which is structurally similar to extracting requirements manually in the first place. The primary benefit of the AI-assisted approach was therefore not efficiency, but coverage consistency: the initial pass covered the entire document systematically, reducing the risk of missing normative statements that a manual skim might overlook.

Beyond the initial extraction, the LLM agent remained a contributor throughout the ongoing maintenance of the verification plan. The MCP-constrained tooling described in Section 3.4 made this safe and efficient: rather than handing an agent write-access to a structured file, each update was channeled through schema-validated tool calls. This allowed the plan to evolve continuously as implementation and verification work progressed - adding traceability links, updating status fields, inserting newly identified requirements - without risking data corruption or requiring manual editing of TOML syntax. The agent's role here was narrowly administrative: executing specific, well-defined updates that the engineer had already decided to make. All judgment calls about what a requirement means, how it should be verified, and whether it is complete remained with the engineer.

4 Design and Architecture

4.1 Ramifications of the Requirements Model on Initial Design Strategy

The structure of the requirements model had direct consequences for the initial design strategy, in ways that were not fully anticipated at the outset.

The **layer dimension** mapped naturally onto the ISO standard's own layered reference model, making a layered module architecture look like the obvious and well-motivated implementation strategy. A dedicated hardware module for each layer - MAC, LLC, fault confinement, and PCS - would allow requirements pertaining to a given layer to be verified in isolation against that layer's module, rather than through the surface of a fully integrated system where internal behavior is obscured by surrounding logic. The

observability dimension reinforced this directly: black-box requirements mapped cleanly onto port-level stimulus and observation, while white-box requirements pointed toward the need for reference models or PSL assertions on internal signals, both of which are most tractable in a per-module testbench. In retrospect, this was a sound conclusion. The modular architecture proved to be the right design choice, and the requirements table provided a well-motivated rationale for it from the start.

The **TX/RX side dimension** had a subtler and more consequential effect. Organizing requirements along the transmitter/receiver axis made intuitive sense from a specification perspective - the ISO standard itself frames many requirements in terms of transmitter behavior and receiver behavior - and it was genuinely useful for thinking through which requirements belonged where. However, it also made a split-path implementation architecture look like the natural design strategy, simply because the requirements were literally organized along that split. The implication appeared to be: implement a TX module, implement an RX module, and map the TX requirements to the former and the RX requirements to the latter.

This turned out to be a red herring. The pitfalls of the split-path approach were not at all apparent from the requirements table alone. The table made the split architecture look clean and well-motivated. The problems - design drift between separately implemented FSMs, integration complexity, and unnecessary hardware duplication - only surfaced later, during integration. This is an important general lesson: the structure of a requirements model can inadvertently bias architectural decisions in ways that are not immediately obvious, and the apparent naturalness of a design strategy that mirrors the requirements structure is not in itself a reliable signal that the strategy is sound.

4.2 Architectural Design Decisions

Before settling on the final architecture, several design alternatives were evaluated. The exploration drew on three sources: the existing in-house CAN Classic controller (Section 1.2), the open-source CTU CAN FD core [2], [3], and the ISO 11898-1 standard's own layered reference model [1]. The protocol requirements and engineering constraints documented in Section ?? bound the feasible design space; this section documents the key decisions made within it.

4.2.1 Monolithic vs. Layered Architecture

The most fundamental architectural decision was whether to extend the existing monolithic controller or adopt a layered decomposition.

The existing controller uses a flat architecture: a single FSM (`can_fsm`, 810 lines) directly drives the bus output, reads the bus input, manages bit counting, handles stuff bit insertion, coordinates CRC computation, and updates error counters - all in one clocked process. This approach minimizes inter-module latency and signal fanout, since every decision is made in a single combinational cone. For CAN Classic, where the frame has at most two format variants (base and extended) and a single bit rate, the complexity is manageable.

CAN FD, however, introduces four frame formats, dual bit rates, fixed bit stuffing with SBC encoding, three CRC polynomials, and a richer error model. Extending the monolithic FSM to cover these features would require nested conditionals on the format type in nearly every state, increasing the cyclomatic complexity well beyond what is practical to verify or maintain. The existing controller's `s_arbitration` state already contains 70 lines of interleaved TX/RX logic for two frame formats; scaling this to four formats with additional control fields (BRS, ESI, SBC) would roughly triple the state's complexity.

The alternative - and the approach adopted - is to follow the ISO 11898-1 reference model, which decom-

poses the data link layer into LLC, MAC, and PCS sub-layers with a separate FCE. Each sub-layer has a well-defined service interface (described in Section ??), and each can be implemented and verified independently. This decomposition introduces inter-module interfaces and pipeline latency, but it confines format-specific complexity to the module where it belongs: the MAC FSM handles frame field sequencing, the bit stuffer handles stuff bit insertion and SBC generation, and the CRC engine handles polynomial computation. No module needs to know about all three concerns simultaneously.

CTU CAN FD [2] takes a middle path: it separates protocol processing (in a CAN Core block) from buffer management and register access, but the protocol core itself remains a large monolithic unit. This is a reasonable trade-off for a general-purpose IP core that must support message filtering, multiple TX buffers, and DMA - features that require tight coupling between the protocol engine and the buffer subsystem. For the narrower scope of this project (protocol engine only, no message RAM or filtering), the full ISO layered decomposition is feasible and preferred because it directly maps each module to a testable subset of the standard's requirements.

4.2.2 Combined vs. Separated TX/RX FSMs

The existing controller uses a single FSM for both transmission and reception, switching between roles via an `is_transmitter` flag. Every state contains parallel `if transmit_i` and `elsif sample_rx_i` branches, one for the transmitter path and one for the receiver path. This approach has two advantages: it avoids duplicating shared state (bit counters, format detection) and it naturally handles the transmitter-to-receiver role switch that occurs when a node loses arbitration.

The disadvantage is that TX and RX logic evolve together. Adding a TX-only feature (such as the data-phase bit rate switch) requires modifying states that also contain RX logic, and vice versa. In CAN FD, the TX and RX paths diverge more than in CAN Classic: the transmitter drives BRS and ESI based on local state, while the receiver samples and validates them. The transmitter feeds the CRC engine with transmitted bits, while the receiver feeds it with received bits (including dynamic stuff bits in the arbitration region, per ISO 6.6.13.3). Encoding both paths in every state produces deeply nested conditionals that are difficult to review and difficult to cover in verification.

The chosen design uses separate TX and RX FSMs (`can_mac_fsm_tx` with 19 states, `can_mac_fsm_rx` with 17 states), each with its own bit stuffer and CRC interface. The two FSMs share no state. Arbitration loss in the TX FSM triggers a `transfer_status` change, but the actual reception of the frame is handled by the RX FSM operating independently on the same bus. This separation means that the TX FSM can be verified exhaustively against TX-side requirements without any RX stimulus beyond bus-level loopback, and conversely for the RX FSM.

The cost of separation is that sub-components used by both paths - the bit stuffer (`can_mac_bs`) and CRC engine (`can_mac_crc`) - must be either shared with multiplexing logic or duplicated. The chosen design duplicates both: `can_mac_tx` instantiates its own `can_mac_bs` and `can_mac_crc`, and `can_mac_rx` does the same. This trades a modest increase in area for complete independence between the two paths - each FSM drives its bit stuffer and CRC engine through its own interface signals without any arbitration or routing logic. The duplication is justified because the bit stuffer and CRC engine are small combinational-plus-register modules (under 120 and 80 lines respectively), while a shared-resource approach would require multiplexing logic in the `can_mac` wrapper and careful coordination to prevent one path's state from corrupting the other's. The `can_mac` wrapper (Section 4.6.3) therefore contains no datapath logic at all - it is purely structural, instantiating `can_mac_tx`, `can_mac_rx`, and `can_fce`, and merging only the FCE error signals.

4.2.3 Per-Field vs. Per-Phase FSM Granularity

A second FSM design decision was the granularity of states. The existing controller uses per-phase states: `s_arbitration` covers all ID bits, SRR, IDE, and RTR; `s_control` covers reserved bits and DLC; `s_data` covers all data bytes. Within each state, a bit counter and conditional logic dispatch the correct action for each bit position.

This per-phase approach keeps the state count low (18 states in the existing controller) but pushes complexity into the bit-counting logic. The `s_arbitration` state must distinguish between base and extended formats using the bit counter, handle the SRR/IDE branching point at bit 12/13, and detect arbitration loss - all in a single state with a single counter.

The alternative is per-field states, where each protocol field (SOF, ID, SRR/RRS, IDE, FDF, RES, BRS, ESI, DLC, Data, SBC, CRC, ACK, EOF) has its own state. This increases the state count (19 TX states, 17 RX states) but eliminates most bit-counter conditionals: when the FSM is in `s_brs`, it knows exactly which bit it is processing without consulting a counter. Format-dependent transitions are handled by the state graph itself - for example, the FSM transitions from `s_ide` to `s_fdf_r1_r0` for all formats, but from `s_fdf_r1_r0` it may go to `s_dlc` (Classic) or `s_res_r0` then `s_brs` (FD). Each state contains only the logic relevant to its specific field, and named guard predicates (`v_in_arbitration_field`, `v_in_dynamic_stuff_field`, `v_in_fixed_stuff_field`) replace repeated bit-range comparisons.

The per-field approach was chosen because it aligns with the verification plan: each field in the frame structure corresponds to a set of requirements in the verification plan, and each FSM state can be traced directly to those requirements. It also simplifies formal verification, since PSL assertions can reference state names rather than counter ranges.

4.2.4 Interface Record Design

The company's `std_logic`-only port constraint does not preclude the use of VHDL record types on entity ports - records of `std_logic` and `std_logic_vector` fields satisfy the constraint. The design uses typed record interfaces (e.g., `t_can_mac_pcs_if_m2s`, `t_can_mac_fsm_bs_if_s2m`) to bundle related signals into a single port. Each record type has a corresponding reset constant (e.g., `c_can_mac_pcs_if_m2s_reset`), ensuring that every module can be reset to a known state without manually enumerating each field.

The alternative - flat port lists with individual `std_logic` signals - was used in the existing controller, where the FSM entity has 20 individual ports for its various sub-module interfaces. This approach becomes unwieldy as the number of inter-module signals grows: the new design has over 60 inter-module signals across its interfaces, and bundling them into records reduces the port list from dozens of lines to six record-typed ports per FSM entity.

The naming convention follows the data-flow direction: `m2s/s2m` (master-to-slave/slave-to-master) for control interfaces, and `s2d/d2s` (source-to-destination/destination-to-source) for Avalon-ST data-transfer interfaces. This convention is consistent with the company's existing interface naming and makes the direction of data flow explicit at every port.

4.2.5 Dual CRC Data Feeds

A non-obvious design decision concerns the CRC engine's data input. In CAN Classic, the CRC is computed over the raw bit stream excluding stuff bits. In CAN FD, the CRC is computed over the raw bit stream in

most of the frame, but in the arbitration region the computation includes dynamic stuff bits - a difference from Classic that exists because the FD CRC must protect the stuff bit count as well as the data. This means that a single data feed to the CRC engine is insufficient: the Classic CRC needs the de-stuffed stream, while the FD CRC needs the stuffed stream during arbitration and the de-stuffed stream elsewhere.

The chosen solution exposes two data inputs on the CRC interface: `data_cc` (Classic CAN data, always de-stuffed) and `data_fd` (FD data, which includes dynamic stuff bits during the arbitration region). The FSM drives both feeds, and the CRC wrapper routes `data_cc` to the CRC-15 engine and `data_fd` to the CRC-17 and CRC-21 engines. This avoids multiplexing logic inside the CRC module and keeps the CRC wrapper purely structural - it instantiates three `gen_crc` blocks and an output mux, with no protocol knowledge.

4.3 System Overview

A complete CAN node decomposes into a TX path and an RX path, coordinated by a shared Fault Confinement Entity (FCE) and Physical Medium Attachment (PMA) control, as shown in Figure 3. Each path spans three sub-layers - LLC (Section 4.5), MAC (Section 4.6), and PCS (Section 4.8) - with the LLC frame format defined in Section 4.4 and interface bundles defined in Section ???. A centralized types package (Section ???) defines all protocol constants and interface records. Within the MAC sub-layer, a unified `can_mac` wrapper (Section 4.6.3) instantiates `can_mac_tx`, `can_mac_rx`, and `can_fce`, merging their error signals internally so that the wrapper exposes only LLC and PCS interfaces for each path plus the FCE's LLC and PCS interfaces.

4.4 LLC Frame Format

The LLC Frame format used by the existing CAN-bus implementation (`can_bus_controller`) is depicted in Figure 4 with the ID byte format depicted in Table 4. A revised LLC Frame supporting FD is depicted in Figure 5. The revised format adds a 3-bit FMT [1, Sec. 6.4.3] field to the DLC byte (LLC Frame byte 4), expands the data field to 64 bytes, and repurposes reserved bits in the last byte for BRS and ESI flags.

The FMT encodes the supported frame formats as '000' = CB, '100' = CE, '010' = FB, '110' = FE. Accordingly, for the frame content to be self-consistent the IDE bit must be set to 0 for FMT = CB/FB and 1 for FMT = CE/FE. The revised format is designed to be backward compatible. An implementation that only supports Classic frames can simply ignore the FMT bits and treat all frames as Classic, while an implementation that supports FD can use the FMT field and the additional control bits (BRS and ESI) to distinguish frame types without affecting the existing ID and data field structure.

Table 4: ID byte encoding for the LLC frame formats shown in Figure 4 and Figure 5.

Format	Byte ID0	Byte ID1	Byte ID2	Byte ID3
Basic	'00000000'	'00000000'	'00000' & 'ID(10-8)'	'ID(7-0)'
Extended	'000' & 'ID(28:24)'	'ID(23-16)'	'ID(15-8)'	'ID(7-0)'

4.5 LLC Sub-layer

Responsible for frame buffering and retransmission management. It provides an Avalon-ST interface to the user application and communicates with the FCE to handle retransmission limits and error status reporting.




Figure 3: CAN node decomposition across LLC, MAC, PCS, FCE, and PMA boundaries. Interface definitions are provided for `llc_tx_if` (Table ??), `llc_rx_if` (Table ??), `llc_mac_tx_if` (Table ??), `llc_mac_rx_if` (Table ??), `mac_pcs_if` (Table ??), `auif_if` (Table ??), `fce_llc_if` (Table ??), `fce_mac_if` (Table ??), and `fce_pcs_if` (Table ??).




Figure 4: Current LLC frame format (15 bytes). ID byte encoding is defined in Table 4.




Figure 5: Revised LLC frame format with FD support, shown at maximum length (71 bytes). The FMT field selects frame type; BRS and ESI repurpose reserved bits in the final byte. ID byte encoding is defined in Table 4.

4.5.1 can_llc_tx

can_llc_tx buffers an incoming LLC frame from the user, streams it byte-by-byte to the MAC serializer, and manages retransmission on bus disturbance up to the ISO-mandated limit [1, Sec. 6.4.5] and 6.5.

4.5.2 can_llc_rx

can_llc_rx receives a reconstructed LLC frame from the MAC layer, applies acceptance filtering, and delivers accepted frames to the user over the llc_rx_if Avalon-ST stream [1, Sec. 6.4.5].

4.6 MAC Sub-layer

The MAC sub-layer is the core of the protocol logic, responsible for bit serialization, CRC generation, bit stuffing, and frame-level error detection. It coordinates closely with the FCE (Section 4.7) for error counter management and node-state transitions (Error Active/Passive/Bus Off), and with the PCS (Section 4.8) for sample-point-driven bit output.

The earlier CAN bus controller concentrated TX, RX, MAC, and FCE logic in a single monolithic FSM (can_fsm), with the sub-functions - serialization (can_ast_to_serial), bit stuffing (can_stuff_bit_gen), and CRC (gen_crc) - implemented in satellite modules driven directly by it. PCS timing logic was implemented in can_node_clock, which fed sample-point and transmit pulses into can_fsm - a strategy retained in the current design.

The current design restructures these modules around the ISO 11898-1 [1] layer boundaries. On the TX side the mapping is direct: can_ast_to_serial becomes can_mac_ser_tx (Section 4.6.1.2), can_stuff_bit_gen becomes can_mac_bs (Section 4.6.1.3), gen_crc becomes can_mac_crc (Section 4.6.1.4), and can_node_clock becomes can_pcs_tx (Section 4.8.1). The bit stuffer and CRC engine are reused across both paths - each of can_mac_tx and can_mac_rx instantiates its own copy of the same can_mac_bs and can_mac_crc entities, with the RX path reusing the bit stuffer for destuffing and the CRC engine for CRC checking. The RX FSM (can_mac_fsm_rx) stores received bits directly in an internal frame array and streams the completed frame to the LLC during the quiet phase, eliminating the need for a separate deserializer module. Fault confinement, previously embedded in can_fsm, is separated into its own entity (can_fce) and wired through the unified can_mac wrapper (Section 4.6.3).

4.6.1 can_mac_tx

The can_mac_tx (Figure 6) is composed of four main components:

1. **can_mac_ser_tx**: Converts LLC bytes into a serial bit stream, extracts LLC metadata from config bytes
2. **can_mac_fsm_tx**: Coordinates frame transmission with per-field state granularity, controls all sub-modules
3. **can_mac_bs**: Shared bit stuffer implementing both dynamic and fixed bit stuffing modes
4. **can_mac_crc**: CRC engine with three parallel generators (CRC-15, CRC-17, CRC-21) and dual data feeds for CC/FD

4.6.1.1 can_mac_fsm_tx can_mac_fsm_tx (Figure 7) orchestrates frame transmission, coordinating the serializer (Section 4.6.1.2), bit stuffer (Section 4.6.1.3), and CRC engine (Section 4.6.1.4). The FSM uses per-field state granularity: rather than a single s_transmitting_frame state with function-dispatched field logic, each protocol field (SOF, ID, DLC, data, CRC, ACK, EOF, etc.) has its own dedicated FSM state.




Figure 6: `can_mac_tx` architecture showing internal component interconnections and external interfaces. Internal interface definitions are provided for `can_mac_ser_fsm_if` (Table ??), `can_mac_fsm_bs_if` (Table ??), and `can_mac_fsm_crc_if` (Table ??). The bit stuffer and CRC engine are the same entities reused in `can_mac_rx`.

This design makes the field boundaries explicit in the state encoding and eliminates the need for frame parameter precomputation functions. The FSM derives `data_len` and `crc_length` from `t_llc_metadata` on frame entry and tracks `bit_count` within each field state.

On each sample-point strobe from `can_pcs_tx` (Section 4.8.1), each frame-field state executes the following sequence:

1. **Monitor the bus.** The sampled bus polarity is compared against the previously transmitted bit to detect errors, ACK, and arbitration loss. In the CAN-FD data phase, the FSM maintains a 32-bit polarity history shift register for Transmitter Delay Compensation (TDC). Any pending secondary sample point (SSP) observation from the previous bit period is evaluated against this history before the current bit is checked [1, Sec. 7.3.4]. A detected error triggers a transition to `s_error_overload` with the appropriate flag type based on the FCE fault confinement status [1, Secs. 8.1.3–8.1.4]. Arbitration loss during the ID or control fields causes the FSM to stop driving and return to `s_intermission`.
2. **Determine the next bit.** If the bit stuffer has a pending stuff bit (`bs_i.valid`), that takes priority. Otherwise, the polarity is determined by the current state: form bits (SOF, IDE, FDF, reserved, delimiters, EOF) have fixed polarities, while ID, DLC, data, SBC, and CRC bits are sourced from the serializer, metadata, bit stuffer count, or CRC register respectively.
3. **Feed the CRC engine and bit stuffer.** The output polarity is forwarded to the CRC engine (for bits within the CRC-protected region) via separate `data_cc` and `data_fd` feeds, and to the bit stuffer (for bits within the stuffing region). The FSM asserts `fixed_bit_stuffing_en` when entering the SBC/CRC field region of FD frames to switch the bit stuffer from dynamic to fixed mode.
4. **Present the bit at the PCS interface.** The resolved polarity is driven onto the `t_can_mac_pcs_if_m2s` interface. The `use_data_rate` signal is asserted during the CAN-FD data phase to switch the PCS to faster bit timing, and `start_tdc` is pulsed at the FDF bit to initiate transmitter delay compensation [1, Sec. 7.3.4].

4.6.1.2 can_mac_ser_tx `can_mac_ser_tx` converts the LLC byte stream into a serial polarity bit stream for the MAC FSM. Its four-state FSM (Figure 8) manages the two-byte configuration handshake, byte fetching, and bit-by-bit serialization. The serializer extracts LLC metadata (IDE, FDF, DLC, FTYP, BRS, ESI) from the two config bytes and registers it in `t_llc_metadata`, which remains stable for the entire frame. The serializer also handles ID field padding - skipping unused bits in the 32-bit ID field for 11-bit base identifiers - and forwards `transfer_status` from the FSM back to the LLC to terminate serialization on error or completion.

4.6.1.3 can_mac_bs `can_mac_bs` implements both dynamic and fixed bit stuffing for CAN Classic and CAN-FD frames [1, Sec. 10.6]. The same entity is instantiated in both `can_mac_tx` and `can_mac_rx` - on the TX side it inserts stuff bits, while on the RX side the FSM uses its output to detect and remove stuff bits from the received stream.

In **dynamic mode** (`fixed_bit_stuffing_en = '0'`), the stuffer monitors the polarity stream and inserts an inverse-polarity stuff bit after every five consecutive identical bits (ISO 6.6.13.2). In **fixed mode** (`fixed_bit_stuffing_en = '1'`), used for the FD CRC region, one fixed stuff bit (FSB) is inserted on the rising edge of `fixed_bit_stuffing_en`, then one every four real bits (ISO 6.6.13.3.1). Any dynamic stuff bit pending at the mode boundary is suppressed. The stuffer maintains a Gray-coded stuff bit count with parity for the SBC field via `f_to_gray()` and `f_calc_parity()`. The data path diagram is shown in




Figure 7: can_mac_fsm_tx FSM with per-field state granularity. Frame-field states (s_{sof} through s_{eof}) each handle one protocol field, with bit_count tracking position within the field. After a successful frame or arbitration loss the FSM passes through $s_{\text{intermission}}$ before returning to $s_{\text{bus_idle}}$. Detected errors branch to $s_{\text{error_overload}}$, which drives either an active (dominant) or passive (recessive) error flag depending on FCE status. The dashed box groups the frame-field states for clarity.




Figure 8: can_mac_ser_tx FSM. The serializer accepts LLC frame bytes over a ready/valid handshake and shifts them out MSB-first to the MAC FSM one bit per ready pulse. LLC metadata is extracted from the two config bytes and registered in t_llc_metadata for use by the FSM throughout the frame.

Figure 9.




Figure 9: `can_mac_bs` data path diagram. When `valid` and `data` = `last_polarity`, `consecutive_count` increments. Otherwise it restarts at 1 with `last_polarity` updated to `data`. When `consecutive_count` reaches `stuff_width` (5 in dynamic mode, 4 in fixed mode), `valid` and inverted `data` are driven and `stuff_count` increments. The `to_gray()` + `calc_parity()` block encodes `stuff_count` into the `stuff_bit_count` output. The external `rst` signal resets `consecutive_count` and `stuff_count` to zero. The `fixed_bit_stuffing_en` signal selects between dynamic and fixed modes. All outputs are registered.

4.6.1.4 can_mac_crc `can_mac_crc` provides CRC generation and checking for both CAN Classic and CAN-FD frames. CAN Classic frames use CRC-15, while CAN-FD frames use CRC-17 (data payloads up to 16 bytes) or CRC-21 (data payloads above 16 bytes) [1, Sec. 10.4.2.6]. The same entity is instantiated in both `can_mac_tx` (for generation) and `can_mac_rx` (for checking). The FSM selects the appropriate polynomial via the `crc_poly_select` field (Table ??). The data path diagram is shown in Figure 10.

Three parallel `gen_crc` instances run continuously on separate data feeds: `data_cc` drives CRC-15 via `valid_cc`, while `data_fd` drives both CRC-17 and CRC-21 via `valid_fd`. This dual-feed architecture is necessary because CC and FD frames compute CRC over different bit streams (CC excludes stuff bits; FD includes them), and the RX path does not know which CRC engine to use until after the frame type has been determined. The output multiplexer selects the active engine's result based on `crc_poly_select` and left-aligns it to the common 21-bit output width.

4.6.2 can_mac_rx

`can_mac_rx` receives the bit stream from the PCS, performs bit destuffing and CRC verification, and re-constructs the LLC frame for delivery to `can_llc_rx` [1, Sec. 6.6]. Unlike the TX path, the RX path has no separate deserializer module. The RX FSM (`can_mac_fsm_rx`) stores received bits directly into an internal `llc_frame` array during reception and streams the completed frame byte-by-byte to the LLC during the quiet phase (after EOF and intermission) using a dedicated Avalon-ST streaming process.

The `can_mac_rx` wrapper instantiates three components:

1. `can_mac_fsm_rx`: Frame reception FSM with 17 per-field states mirroring the TX FSM structure
2. `can_mac_bs`: Shared bit stuffer entity, reused for destuffing
3. `can_mac_crc`: Shared CRC engine entity, reused for CRC checking

The RX FSM (Figure 11) validates SBC and CRC fields during reception, detects form errors (reserved bits,




Figure 10: `can_mac_crc` data path diagram. Three parallel `gen_crc` instances accumulate continuously: `data_cc` drives CRC-15 via `valid_cc`, while `data_fd` drives both CRC-17 and CRC-21 via `valid_fd`. The output register selects the active engine result based on `poly_select` and left-aligns it to 21 bits. The external `rst` signal reinitializes all three engines and the output register.

CRC mismatch, delimiter errors), and drives ACK dominant during the ACK slot (1 bit for CC, 2 bits for FD). It reports errors to the FCE through the same `t_can_mac_fce_if_m2s` interface used by the TX FSM.

4.6.3 `can_mac` Unified Wrapper

`can_mac` is a structural wrapper that instantiates `can_mac_tx`, `can_mac_rx`, and `can_fce` (Section 4.7), wiring the FCE internally so that the wrapper exposes only LLC and PCS interfaces for each path plus the FCE's LLC and PCS interfaces. TX and RX have separate PCS interfaces (half-duplex; only one path drives the bus at a time). The TX and RX FCE output records (`t_can_mac_fce_if_m2s`) are merged by OR-ing all fields before feeding the FCE input - the half-duplex guarantee ensures no simultaneous assertions from both paths. The FCE's response (`t_can_mac_fce_if_s2m`, carrying `error_passive_request` and `error_active_request`) is fanned back to both TX and RX paths.

4.7 FCE Sub-layer

The Fault Confinement Entity (`can_fce`) implements the error state machine and counter management specified in [1, Secs. 8.1.3–8.1.4]. It maintains two error counters - TEC (Transmitter Error Counter, 0-511) and REC (Receiver Error Counter, 0-255) - and transitions between three states: `s_error_active` (normal operation), `s_error_passive` (TEC or REC > 127), and `s_bus_off` (TEC > 255).

Counter updates follow the ISO 8.1.4.2 rules: TEC increments by 8 on TX errors (unless `counters_unchanged` is asserted), decrements by 1 on successful TX. REC increments by 1 on RX errors during non-error-flag phases, by 8 on primary errors or error-flag-phase errors, and decrements by 1 or clamps to 127 on successful RX. The FCE state machine (Figure 12) transitions between error-active, error-passive, and bus-off based on counter thresholds. Bus-off recovery requires counting 128 idle conditions (11 consecutive recessive bits each) from the PCS via the `idle_condition` signal, or a `normal_mode_request` from the LLC.

4.8 PCS Sub-layer

Handles bit timing and synchronization. It generates the sample point (SP) and secondary sample point (SSP) strobes. It provides bit-level monitoring data to the FCE to detect synchronization and timing errors.

4.8.1 `can_pcs_tx`

`can_pcs_tx` handles bit timing and Transmitter Delay Compensation (TDC) for the TX path [1, Secs. 7.2–7.3]. It generates the sample point (SP) strobe used by the MAC FSM to advance frame state, and - when TDC is configured - a secondary sample point (SSP) strobe for data-phase bit monitoring. The FSM (Figure 13) has four states reflecting the two bit rates and the TDC measurement window. The `measuring_delay` state determines the Transmitter Delay Compensation Value (TDCV, [1, Sec. 7.3.4]) by observing the TX-to-RX propagation delay, which together with the configured TDCO offset defines the SSP position for subsequent data-phase bits.

4.8.2 `can_pcs_rx`

`can_pcs_rx` implements bit timing and synchronization for the RX path. It performs hard and soft synchronization on the incoming bus signal and generates the sample point strobe used by the MAC RX layer to latch each received bit [1, Secs. 7.2–7.3].




Figure 11: can_mac_fsm_rx FSM with per-field state granularity. The RX FSM mirrors the TX FSM structure (Figure 7) but replaces transmission logic with reception, storage, and validation. On detecting a dominant SOF in s_idle, the FSM enters s_id and stores received bits into an internal llc_frame array. After EOF, the FSM transitions through s_intermission and a dedicated streaming process transfers the stored frame to the LLC. Form errors (reserved bit violations, CRC mismatch, delimiter errors) and SBC mismatches trigger s_error_overload. The dashed box groups the frame-field states for clarity.




Figure 12: can_fce FSM (ISO 8.1.4.1, Figure 43). The s_error_active state is the normal operating mode. When either TEC or REC exceeds 127, the FCE transitions to s_error_passive and asserts error_passive_request to both MAC paths. If TEC exceeds 255, the node enters s_bus_off: the FCE issues bus_off_request to the PCS and bus_off to the LLC. Recovery from bus-off requires 128 idle conditions (each 11 consecutive recessive bits) counted from the PCS, or a normal_mode_request from the LLC, after which the counters are reset and the FCE returns to s_error_active.




Figure 13: **can_pcs_tx** FSM. The **measuring_delay** state is entered on the FDF sample point to measure TDCV (Transmitter Delay Compensation Value, [\[1\]](#), sec. 7.3.4). The BRS bit boundary determines whether data-phase timing is used. All non-idle states return to idle when the frame becomes inactive.

5 Implementation

5.1 Type Safety and Packages

The implementation uses custom record types (defined in `can_types_pkg.vhd`) to ensure clean interfaces between modules.

5.2 Bit Timing and TDC

Detailed description of how `tx_pcs` measures propagation delay and calculates the SSP.

5.3 CRC and Bit Stuffing

Implementation details of the flexible CRC generator and the hybrid bit stuffer.

6 Verification and Results




Figure 14: REQ-009.




Figure 15: REQ-010.

Note: This section is currently being populated as the verification plan is executed.

6.1 Testbench Results Summary

Table 5: Testbench execution status and functional coverage summary.

Testbench	Status	Coverage
tx_pcs_tb	Pass	95%
tx_can_tb	Pass	88%
tx_can_protocol_tb	Pass	92%

7 Discussion

Comparison of the implemented architecture against theoretical models. Performance analysis in high-load scenarios.

7.1 Future Work

1. Make the the CRC and BS modules to save area on the FPGA.
2. CAN XL implementation...

8 Conclusion

Summary of work completed and how objectives were met.

9 References

A Verification Plan

Both tables are regenerated automatically from `verification_plan/verification_plan.toml` on each PDF build. The ID field is the join key between them. See Section 3.3 for the meaning of each field.

A.1 Extracted Requirements

The requirements table contains four fields: the identifier, the ISO 11898-1 source clause, the priority assignment, and the paraphrase. Placing priority here allows each assignment to be evaluated directly against the requirement text. The table corresponds to the requirements extraction and prioritisation phases described in Section 3.2 and Section 3.2.5.

Table 6: ISO 11898-1 extracted requirements.

ID	ISO ref	Priority	Requirement
REQ-001	6.4.5.2	P1	LLC shall notify the LLC user of every successful DF transmission or reception (all formats), and every RF transmission or reception (CB, CE only).
REQ-002	6.4.5.5.2	P2	Any transmission request shall be processed no later than the second SOF after the request when no EFs are present.
REQ-003	6.4.5.5.3	P2	Abort request rules: 1) Frames not yet passed to the MAC are aborted immediately. 2) Frames already at the MAC are only aborted on MAC error or arbitration loss. 3) The abort is processed before the second SOF (no EFs) or third SOF (with EFs).
REQ-004	6.5.2, 6.5.3	P1	Retransmission policy: 1) Frames losing arbitration shall be retransmitted up to the configured limit unless aborted. 2) Retransmission attempts shall be counted in the retransmission counter.
REQ-005	6.5.4	P3	Frame acceptance filtering: 1) Receivers should use frame acceptance filtering to decide frame relevance. 2) Filtering should be a single, self-contained operation independent of previous transactions.
REQ-006	6.5.6	P1	LLC shall report five transfer status values: Ongoing, Transmitted, Lost_Arbitration, Disturbed, and Aborted, each reflecting the corresponding MAC sub-layer event.
REQ-007	6.6.4.4	P1	CRC rules: 1) Polynomial selection: CRC_15 for CC, CRC_17 for FD ≤ 16 B payload, CRC_21 for FD > 16 B payload. 2) Init vector: all-zero for CRC_15, MSB-one for CRC_17 and CRC_21.
REQ-008	6.6.5.3, 6.6.6.3	P1	Error/overload delimiter: 1) Error and overload delimiters each consist of eight recessive bits. 2) After the flag, all nodes monitor the bus and start their seven-bit delimiter simultaneously on detecting the first recessive bit.
REQ-009	6.6.7.2	P1	Intermission rules: 1) Intermission shall be exactly three recessive bits. 2) No DF or RF may be started during intermission. 3) A dominant bit at the third intermission bit shall be interpreted as SOF.

ID	ISO ref	Priority	Requirement
REQ-010	6.6.7.3	P1	Bus idle is recognised at: 1) Third recessive bit of intermission - error-active nodes and receivers. 2) Eighth recessive bit of suspend transmission - error-passive transmitters. 3) Upon leaving the bus-integration state - any node.
REQ-011	6.6.7.4	P2	Suspend transmission: 1) An error-passive transmitter shall suspend for 8 bit times after intermission. 2) If another node starts transmitting during the suspend window, the node shall become a receiver.
REQ-012	6.6.7.5	P1	Bus re-integration: 1) Nodes shall start bus re-integration on protocol startup or bus-off recovery. 2) Re-integration completes after 11 consecutive recessive bits at the sample point.
REQ-013	6.6.8	P1	Frame initiation: 1) A node shall send SOF only when the bus is idle. 2) A dominant third intermission bit causes a node with a pending frame (if error-active or previous receiver) to skip the SOF and begin arbitration immediately.
REQ-014	6.6.10.1, 6.6.10.4	P2	Remote frame (CB, CE): 1) RTR bit shall be dominant for data frames and recessive for remote frames. 2) Remote frames have no data field; DLC indicates the requested data length.
REQ-015	6.6.10.2, 6.6.11.2	P2	SRR and RRS reserved bits: 1) SRR (CE, FE): transmitted recessive; receivers accept dominant or recessive without flagging a form error. 2) RRS (FB, FE): transmitted dominant; receivers accept recessive or dominant without flagging a form error.
REQ-016	6.6.10.5, 6.6.11.5	P1	CRC bit stream coverage: 1) CC: SOF, arbitration, control, data fields (dynamic stuff bits excluded). 2) FD: SOF, arbitration, control, data, dynamic stuff bits, and the SBC field (fixed stuff bits excluded).
REQ-017	6.6.10.6, 6.6.11.6	P1	ACK slot rules: 1) Receivers shall ACK consistent frames and flag inconsistent frames with an EF. 2) Transmitters shall flag unacknowledged frames with an EF. 3) FD frames tolerate a two-bit dominant ACK overlap to compensate for receiver phase differences.

ID	ISO ref	Priority	Requirement
REQ-018	6.6.11.3	P1	res bit (FB, FE): 1) The res bit shall be dominant. 2) A recessive res bit is a form error. 3) The bit rate switches to the FD data rate at the res bit.
REQ-019	6.6.11.3	P2	ESI bit encoding: 1) Recessive when the LLC ESI flag is set. 2) Dominant when the LLC ESI flag is not set and the node is error-active. 3) Recessive when the LLC ESI flag is not set and the node is error-passive.
REQ-020	6.6.11.5	P1	The SBC field shall be Gray-coded with a parity bit and placed before the FCRC, encoding the count of dynamic stuff bits in the frame.
REQ-021	6.6.11.5	P2	The transmitter shall send one recessive CRC delimiter bit; the 2-bit FD ACK window provides tolerance for a second recessive bit before the ACK slot, compensating for receiver phase differences.
REQ-022	6.6.13.2	P1	Dynamic bit stuffing: 1) After five consecutive identical bits, a complementary dynamic stuff bit shall be inserted. 2) CC stuffing covers SOF through the FCRC sequence. 3) FD stuffing covers SOF through the data field only.
REQ-023	6.6.13.3.1	P1	Fixed stuff bits (FB, FE): FD frames shall have a fixed stuff bit before the SBC field and after every fourth CRC bit, each the inverse of its preceding bit.
REQ-024	6.6.15.2	P1	EOF validation: 1) A receiver shall validate a frame after the sixth EOF bit. 2) A dominant seventh EOF bit triggers an OF rather than a form error. 3) A transmitter requires all seven EOF bits to be error-free.
REQ-025	6.6.17.4	P1	Arbitration rules: 1) Recessive-sent, dominant-monitored: lose arbitration and become a receiver. 2) Dominant-sent, recessive-monitored: bit error. 3) Arbitration spans from the first ID bit to the format-specific boundary bit (IDE for CB, RTR for CE, bit after FDF for FB/FE).

ID	ISO ref	Priority	Requirement
REQ-026	6.6.21.2	P1	Bit error rules: 1) A bit error shall be detected on any sent/monitored mismatch. 2) Exception: recessive non-stuff bits during arbitration or the ACK slot. 3) Exception: dominant bits monitored while sending a passive error flag.
REQ-027	6.6.21.2	P1	Stuff error rules: a stuff error shall be detected at the sixth consecutive identical bit in any dynamically-stuffed field.
REQ-028	6.6.21.2	P1	Form error rules: 1) A form error shall be detected on any unexpected fixed-form bit value or wrong fixed stuff bit polarity (FB, FE). 2) Exception: dominant bits at the last EOF bit or at the last bit of any error or overload delimiter.
REQ-029	6.6.21.2	P1	CRC and ACK error rules: 1) Receivers shall compute CRC identically to the transmitter and detect a CRC error on FCRC mismatch. 2) FD frames additionally check SBC match and SBC parity. 3) A transmitter shall detect an ACK error if no dominant bit is monitored in the ACK slot.
REQ-030	6.6.21.3.1	P1	Error flag initiation: 1) On any detected error, start an EF at the immediately following bit and inform the LLC. 2) Data-phase errors require a switch back to nominal bit rate before the EF starts.
REQ-031	6.6.21.3.1	P2	On an FCRC error, a receiver shall: 1) Withhold ACK. 2) Send an EF starting at the first EOF bit (CC frames). 3) Send an EF three bit-times after the CRC delimiter (FD frames).
REQ-032	6.6.21.3.2	P2	An OF is triggered by: 1) LLC request - starts at the first bit of intermission. Implementation is optional. 2) Reactive - dominant at either of the first two intermission bits, the last EOF bit (receivers only), or the last error/overload delimiter bit. The OF starts one bit after the triggering bit.
REQ-033	7.3.2	P1	Nodes shall support separate nominal and FD data bit rate configurations (XL excluded), each based on a programmable integer prescaler that derives the time quantum from the node clock period.

ID	ISO ref	Priority	Requirement
REQ-034	7.3.2	P2	Propagation segment budget: 1) prop_seg shall cover the full round-trip bus delay. 2) Formula 2 requires $\text{prop_seg} \geq t_{\text{node}}(A) + t_{\text{node}}(B) + 2 \times t_{\text{busline}}$, where each node delay includes transceiver round-trip delay.
REQ-035	7.3.4, 6.6.21.3.1	P1	Transmitter Delay Compensation (TDC): 1) FD-capable nodes shall implement TDC for the FD data phase when BRS is recessive, with programmable enable and prescaler constrained to 1 or 2. 2) SSP position = measured transmitter delay + configured offset. 3) When TDC is active, bit errors at the regular SP are ignored; errors detected at the SSP are reacted upon at the following nominal SP.
REQ-036	7.3.5.1	P1	Synchronisation rules: 1) Only one synchronisation is allowed per bit time; synchronisation is re-enabled after the next recessive sample point. 2) Edges qualify only if the previous sample point was recessive and no positive phase error is present while sending dominant. 3) Hard synchronisation occurs on IFS edges, during bus-integration, and at the FDF-to-res transition. 4) All other qualifying edges cause resynchronisation, except during the FD data phase.
REQ-037	7.3.5.3, 7.3.5.4	P1	Synchronisation corrections: 1) Hard synchronisation restarts the bit time with Sync_Seg completed, unconstrained by SJW. 2) Resynchronisation adjusts Phase_Seg1/Phase_Seg2 by SJW when phase error exceeds SJW, otherwise acts identically to hard synchronisation.
REQ-038	7.3.6	P2	Oscillator frequency tolerance: 1) Tolerance df shall satisfy $df < \text{SJW} / (20 \times \text{tbit})$ and $df < \min(\text{Phase_Seg1}, \text{Phase_Seg2}) / (2 \times (13 \times \text{tbit} - \text{Phase_Seg2}))$. 2) SJW shall not exceed $\min(\text{Phase_Seg1}, \text{Phase_Seg2})$.
REQ-039	6.6.7.5, 8.1.4.4	P1	PCS bus-off behaviour: 1) During bus_off, the PCS shall not drive dominant bits. 2) The PCS counts consecutive recessive bits at the sample point and strobes idle_condition after 11, resetting on any dominant.

ID	ISO ref	Priority	Requirement
REQ-040	7.4.2.2, 7.4.2.3	P1	PCS bus-off signalling: 1) The PCS shall disconnect the node from the bus on a bus_off_request from the FCE supervisor. 2) The PCS shall reconnect it on a bus_off_release_request.
REQ-041	8.1.3.2, Table 14	P2	Normal_mode_request resets TEC and REC to zero.
REQ-042	8.1.4.2 rule a), rule b), rule e), rule h)	P1	REC rules: 1) +1 on any receiver-detected error (except bit errors during error/overload flag). 2) +8 on first dominant bit after sending an error flag. 3) +8 on bit error while sending an active error or overload flag. 4) -1 on successful RX if in 1-127, stays 0 if already 0, or is set to 119-127 if above 127.
REQ-043	8.1.4.2 rule c), rule d), rule g)	P1	TEC rules: 1) +8 when sending an error flag (exceptions: error-passive ACK error, pre-arbitration stuff error). 2) +8 on bit error while sending an active error or overload flag. 3) -1 on successful TX (floor 0).
REQ-044	8.1.4.2 rule f)	P2	Prolonged dominant sequence: 1) Counters increment by 8 after 14 consecutive dominant bits following an active error or overload flag (8 after a passive error flag). 2) A further +8 applies for each additional 8 consecutive dominant bits.
REQ-045	8.1.4.3, 8.1.4.4	P1	Error confinement state transitions: 1) Either counter exceeding 127 causes error-passive. 2) Both counters at or below 127 restores error-active. 3) TEC exceeding 255 causes bus-off. 4) Bus-off recovery requires 128 idle conditions with both counters reset to zero.

A.2 Verification Plan Fields

The verification plan table contains the verification metadata fields added during the planning phase. Cross-reference to the requirements table above via the shared ID field.

Table 7: ISO 11898-1 verification plan.

ID	Layer	Side	Format	Observability	Method	Status	Label	File
REQ-001	LLC	both	CB, CE, FB, FE	black_box	simulation	not_started		
REQ-002	LLC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started		
REQ-003	LLC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started		
REQ-004	LLC	transmitter	CB, CE, FB, FE	white_box	code_inspection	not_started		
REQ-005	LLC	receiver	CB, CE, FB, FE	white_box	code_inspection	not_started		
REQ-006	LLC	transmitter	CB, CE, FB, FE	black_box	simulation	not_started		
REQ-007	MAC	both	CB, CE, FB, FE	white_box	simulation, coverage	in_progress	p_output_vc, p_reset_checker	can_mac_crc_tb.vhd
REQ-008	system	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm::~~900, test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-009	MAC	both	CB, CE, FB, FE	white_box	waveform_inspection	in_progress	test_normal	can_mac_pcs_fce_tb.vhd
REQ-010	MAC	both	CB, CE, FB, FE	white_box	waveform_inspection	in_progress	test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-011	MAC	transmitter	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm::~~486, test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-012	system	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm::~~433, test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-013	MAC	transmitter	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm::~~451, test_normal	can_mac_pcs_fce_tb.vhd
REQ-014	MAC	both	CB, CE	black_box	simulation	not_started		

ID	Layer	Side	Format	Observability	Method	Status	Label	File
REQ-015	MAC	both	CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~546, p_fsm:~548, test_normal	can_mac_pcs_fce_tb.vhd
REQ-016	MAC	both	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm:~299, p_fsm:~216	can_mac_fsm.vhd
REQ-017	system	both	CB, CE, FB, FE	white_box	simulation, waveform_inspection	in_progress	test_normal	can_mac_pcs_fce_tb.vhd
REQ-018	MAC	both	FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~613, test_normal	can_mac_pcs_fce_tb.vhd
REQ-019	MAC	transmitter	FB, FE	white_box	waveform_inspection	in_progress	test_normal	can_mac_pcs_fce_tb.vhd
REQ-020	MAC	both	FB, FE	black_box	simulation	in_progress	p_sbc_checker	can_mac_bs_tb.vhd
REQ-021	MAC	transmitter	FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~768, p_fsm:~796, test_normal	can_mac_pcs_fce_tb.vhd
REQ-022	MAC	both	CB, CE, FB, FE	white_box	simulation, coverage	in_progress	p_stuff_bit_checker	can_mac_bs_tb.vhd
REQ-023	MAC	both	FB, FE	white_box	simulation, coverage	in_progress	p_fsb_checker	can_mac_bs_tb.vhd
REQ-024	MAC	both	CB, CE, FB, FE	white_box	simulation, coverage	in_progress	test_normal	can_mac_pcs_fce_tb.vhd
REQ-025	system	transmitter	CB, CE, FB, FE	black_box	simulation	in_progress	test_lost_arb	can_mac_pcs_fce_tb.vhd
REQ-026	MAC	transmitter	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~283, test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-027	MAC	both	CB, CE, FB, FE	black_box	code_inspection	in_progress	p_fsm:~336	can_mac_fsm.vhd
REQ-028	MAC	both	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm:~343	can_mac_fsm.vhd
REQ-029	MAC	both	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm:~415	can_mac_fsm.vhd
REQ-030	MAC	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~294, p_fsm:~283, test_bus_off	can_mac_pcs_fce_tb.vhd

ID	Layer	Side	Format	Observability	Method	Status	Label	File
REQ-031	MAC	receiver	CB, CE, FB, FE	white_box	code_inspection	in_progress	p_fsm:~357	can_mac_fsm.vhd
REQ-032	MAC	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_fsm:~370	can_mac_fsm.vhd
REQ-033	PCS	both	CB, CE, FB, FE	white_box	simulation	in_progress	test_normal	can_pcs_tb.vhd
REQ-034	system	both	CB, CE, FB, FE	black_box	simulation, coverage	in_progress	test_delay_sweep	can_mac_pcs_fce_tb.vhd
REQ-035	PCS	transmitter	FB, FE	black_box	simulation	in_progress	p_check_tdc_delay	can_pcs_tb.vhd
REQ-036	PCS	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_can_pcs:~177, test_normal	can_pcs_tb.vhd
REQ-037	PCS	both	CB, CE, FB, FE	white_box	code_inspection, waveform_inspection	in_progress	p_can_pcs:~228, test_normal	can_pcs_tb.vhd
REQ-038	PCS	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_normal	can_pcs_tb.vhd
REQ-039	PCS	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_bus_off	can_pcs_tb.vhd
REQ-040	system	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_bus_off	can_mac_pcs_fce_tb.vhd
REQ-041	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_reset, test_bus_off_recovery	can_fce_tb.vhd
REQ-042	FCE	receiver	CB, CE, FB, FE	black_box	simulation	in_progress	test_rule_a, test_rule_b, test_rule_e, test_rule_h	can_fce_tb.vhd
REQ-043	FCE	transmitter	CB, CE, FB, FE	black_box	simulation	in_progress	test_rule_c, test_rule_c_exception, test_rule_d, test_rule_g	can_fce_tb.vhd
REQ-044	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_rule_f_tx, test_rule_f_rx	can_fce_tb.vhd

REQ-045	FCE	both	CB, CE, FB, FE	black_box	simulation	in_progress	test_bus_off, test_bus_off_recovery	can_fce_tb.vhd
---------	-----	------	-------------------	-----------	------------	-------------	--	----------------

- [1] International Organization for Standardization, *Road vehicles — controller area network (CAN) — Part 1: Data link layer and physical coding sublayer*. 2024.
- [2] M. Jeřábek and P. Píša, “CTU CAN FD — open-source CAN FD IP core.” 2019. Available: https://gitlab.fel.cvut.cz/canbus/ctucanfd_ip_core
- [3] M. Jeřábek, “Open-source and open-hardware CAN FD protocol support,” PhD thesis, Czech Technical University in Prague, 2019. Available: <https://dspace.cvut.cz/handle/10467/82630>
- [4] International Organization for Standardization, *Road vehicles — controller area network (CAN) conformance test plan — Part 1: Data link layer and physical signalling*. 2016.
- [5] OpenCores Contributors, “CAN protocol controller — OpenCores.” 2003. Available: <https://opencores.org/projects/can>
- [6] S. V. Nøsen, “Canola — CAN controller in VHDL.” 2020. Available: <https://github.com/svnesbo/canola>
- [7] Robert Bosch GmbH, “M_CAN — CAN FD protocol controller IP module.” 2024. Available: <https://www.bosch-semiconductors.com/ip-modules/can-fd/>
- [8] F. Hartwich, “CAN with flexible data-rate,” in *Proceedings of the 13th international CAN conference (iCC 2012)*, 2012.
- [9] AMD (Xilinx), “CAN FD controller — Vivado design suite product guide (PG223).” 2024. Available: <https://docs.amd.com/r/en-US/pg223-canfd>
- [10] CAST, Inc., “CAN-FD protocol controller.” 2024. Available: <https://www.cast-inc.com/automotive/can/can-fd-protocol-controller>
- [11] Intel Corporation, “Avalon interface specifications,” Intel Corporation, 683091, 2021. Available: <https://www.intel.com/programmable/technical-pdfs/683091.pdf>
- [12] J. Bergeron, “Verification planning,” in *Writing testbenches: Functional verification of HDL models*, 2nd ed., Kluwer Academic Publishers, 2003.